

Preference Discovery in University Admissions: The Case for Dynamic Multioffer Mechanisms

Julien Grenet

Centre National de la Recherche Scientifique and Paris School of Economics

YingHua He

Rice University and Toulouse School of Economics

Dorothea Kübler

WZB Berlin Social Science Center and Technische Universität Berlin

We document quasi-experimental evidence against the common assumption in the matching literature that agents have full information on their own preferences. In Germany's university admissions, the first stages of the Gale-Shapley algorithm are implemented in real time, allowing for multiple offers per student. We demonstrate that non-exploding early offers are accepted more often than later offers, despite not being more desirable. These results, together with survey evidence and a theoretical model, are consistent with students' costly discovery of preferences. A novel dynamic multioffer mechanism that batches early offers improves matching efficiency by informing students of offer availability before preference discovery.

WZB = Wissenschaftszentrum Berlin für Sozialforschung. For constructive comments, we thank Inácio Bó, Caterina Calsamiglia, Estelle Cantillon, Laura Doval, Tobias Gamp, Rustamdjan Hakimov, and Thilo Klein, as well as seminar and conference participants at the Center for Research in Economics and Statistics/École Polytechnique, Econometric Society/Bocconi University Virtual World Congress 2020, Matching in Practice (2017 Brussels), National Bureau of Economic Research Market Design Working Group 2021, the University of Gothenburg, Stanford University, and Zentrum für Europäische Wirtschaftsforschung Mannheim. Jennifer Rontganger provided excellent copyediting service. We are grateful

Electronically published April 6, 2022

Journal of Political Economy, volume 130, number 6, June 2022.

© 2022 The University of Chicago. This work is licensed under a Creative Commons Attribution-NonCommercial 4.0 International License (CC BY-NC 4.0), which permits non-commercial reuse of the work with attribution. For commercial use, contact journalpermissions@press.uchicago.edu. Published by The University of Chicago Press.
<https://doi.org/10.1086/718983>

Research on the design of matching markets has been a success story, not least because it has resulted in improved designs for school choice, university admissions, and entry-level labor markets (Roth and Peranson 1999; Abdulkadiroğlu, Pathak, and Roth 2005; Abdulkadiroğlu et al. 2005). An important contributor to these successes has been market centralization, which aims at promoting efficiency—for instance, by reducing congestion and market unraveling (Roth 1990).

Coupled with growing market centralization is the near-universal adoption of single-offer mechanisms, especially the deferred acceptance (DA) mechanism (Gale and Shapley 1962). Such mechanisms require every agent to rank her potential match partners before at most one match offer per agent is generated.¹ The widespread use of single-offer mechanisms is a direct consequence of the common assumption in the matching literature that every agent has full information on her own preferences (e.g., Roth and Sotomayor 1990; Abdulkadiroğlu and Sönmez 2003). Under this assumption, an agent can effortlessly rank match partners according to her preferences even though she is uncertain about eventual offers.

Yet full information on one's own preferences is unlikely to hold in practice. Take university admissions as an example. A student often has access to hundreds if not thousands of university programs.² Each program is a combination of a university and a major and has multidimensional characteristics, including academic quality, courses offered, and quality of life. A student needs to invest time and effort to learn about program quality and discover her preferences. As a result, it could be welfare-improving to adopt mechanisms that allow students to hold multiple offers before they have to rank programs. In such a dynamic multioffer mechanism, early offers inform students of their admission chances and hence allow them to learn about programs more efficiently.

The first contribution of our study is to provide unambiguous empirical evidence against the assumption that students have full information

to the Stiftung für Hochschulzulassung—in particular, Matthias Bode—for invaluable assistance. Financial support from the National Science Foundation through grant SES-1730636, the French National Research Agency through project ANR-14-FRAL-0005 FDA and EUR grant ANR-17-EURE-0001, and the German Research Foundation through projects KU 1971/3-1 and CRC TRR 190 is gratefully acknowledged. A previous version of this paper was circulated under the title “Decentralizing Centralized Matching Markets: Implications from Early Offers in University Admissions.” Replication files are provided as supplementary material online. This paper was edited by Emir Kamenica.

¹ When a single-offer mechanism is used in many-to-one matching, such as school choice and college admissions, a school that can accept multiple students will receive multiple match offers. However, a student can attend only one school and will obtain at most one offer.

² As more and more market segments are integrated into centralized markets with a single-offer mechanism, the number of potential match partners for a student can be huge. This trend is observed beyond university admissions. For instance, charter school and traditional public school admissions have been unified in a single-offer centralized design in Denver (Abdulkadiroğlu et al. 2017) and New Orleans (Abdulkadiroğlu et al. 2020).

on their own preferences. In an administrative data set on university admissions in Germany, we identify a quasi experiment in which the arrival time of admission offers is as good as random. We show that a student is more likely to accept an early offer relative to later offers, despite the fact that offers do not expire until the end of the admissions procedure. This early-offer effect cannot be reconciled with the assumption that students have full information on their own preferences but is instead consistent with students discovering them at a cost. This interpretation is further corroborated by direct evidence from a survey of students.

Our second contribution is to make a case for dynamic multioffer mechanisms in settings with nonnegligible costs of preference discovery. An example of such a mechanism is used in Germany's university admissions. A student who has received an early offer from a program is certain that this program is available to her and can thus target her learning more efficiently. However, the random arrival of offers may lead students to suboptimally accept early offers. We therefore propose a new multioffer mechanism that batches early offers. The benefits of the new mechanism are shown in a theoretical model and in simulations based on the data from Germany. In particular, the new mechanism dominates the DA, the most widely used single-offer mechanism.

Our empirical analysis exploits a procedure called *Dialogorientiertes Serviceverfahren* (DoSV; literally, dialogue-oriented service procedure) that is used for admissions to overdemanded university programs in Germany. Admission decisions are made by each program separately. The procedure is based on the program-proposing Gale-Shapley (GS) algorithm but differs from the standard implementation in that each student can defer her commitment to a rank-order list (ROL) of programs. Specifically, the DoSV implements the first stages of the GS algorithm in real time and allows students to hold multiple offers. During this phase of 34 days, all offers and acceptances are communicated via a clearinghouse, with students and programs interacting as if in a decentralized market. Programs make admission offers to their preferred students in real time, no offer expires until the end of the procedure, and students can decide to accept an offer and exit the procedure, to retain all offers, or to keep only a subset of their offers, also in real time.

At the end of the multioffer phase, students who have not yet accepted an offer are required to finalize their ROLs and enter a single-offer phase in which the computerized program-proposing GS algorithm is run for the remaining students and seats.³ For a student with offers from the multioffer

³ The sequential filling of seats is also a feature of a separate procedure that was used for admissions to medical programs in Germany (Braun, Dwenger, and Kübler 2010; Westkamp 2013; Braun et al. 2014). There, quotas are filled one after the other, each quota with its own ROLs and matching algorithm (Boston immediate acceptance or GS). In contrast, the DoSV is based on one algorithm (GS) for all seats, but its early stages resemble a decentralized market.

phase, the highest-ranked offer in her final ROL is kept and will not be taken away from her unless she receives an offer from a program that is ranked higher in her final ROL. The single-offer phase ensures that a stable matching (with respect to stated preferences) is reached. Thus, unraveling does not play a role in this market.⁴

Our data set contains every event during the admission process for 2015–16 as well as its exact timing. We define a program as being *feasible* to a student if the student has applied to the program and would receive its admission offer as long as she does not exit the procedure early. There are 21,711 students in our data set who have at least two feasible programs and have accepted one of them. We find that, relative to an offer arriving in the later, single-offer phase, an offer that arrives during the multioffer phase—which we define as an *early offer*—is more likely to be accepted.

The early-offer effect is sizable when we calculate its impact on offer acceptance. In our sample, the average probability of a feasible program being accepted is 0.385. An early offer that is not the first one raises the acceptance probability by 8.7 percentage points—a 27.9% increase. More significantly, the first early offer increases the acceptance probability by 11.8 percentage points—a 38.3% jump.⁵

We rule out a range of possible explanations of the early-offer effect through additional data analyses. The effect is unlikely to be driven by students' preferences for being ranked high by a program, an immediate reaction, pessimism about future offers, the need for a head start in house hunting, or dislike of being assigned by a computerized algorithm.

Our preferred explanation is that the early-offer effect is driven by students' costly discovery of their preferences over programs. From a student survey, we find direct evidence that at the start of the procedure, many students do not yet have clear preferences over programs. Furthermore, they tend to invest more time learning about the universities that made them an early offer, and early offers influence their perceptions of the programs.

These empirical findings can be formalized in a model of university admissions with preference discovery and are illustrated in a numerical example. A student can learn her valuation of a university program only after paying a cost. She applies to two programs and has access to an outside option of a known value. In expectation, each program makes her an offer with some probability, which determines the student's endogenous

⁴ Decentralized markets without such a phase guaranteeing a stable outcome have been studied theoretically by Niederle and Yariv (2009) and experimentally by Echenique and Yariv (2013).

⁵ Another way to gauge the magnitude of the effect is to use distance from a student's home to the university where the program is located as a numeraire. At the sample mean of 126 km, an early offer amounts to reducing the distance to the corresponding program by 61 km. The first early offer has an even larger effect, equivalent to a reduction of 79 km.

learning decision. An early offer from a program, which leads the student to update the corresponding offer probability to one, increases her incentive to learn about this program. Finally, learning about a program increases the probability that its offer is accepted, leading to the early-offer effect.

These results provide novel insights for market design regarding the benefits of dynamic multioffer mechanisms. In a generalized model with preference discovery, we evaluate three mechanisms: (i) the DA, which requires students to finalize their ROLs before receiving any offer; (ii) the multioffer DA (M-DA), of which the DoSV procedure is an example; and (iii) our proposed batched multioffer DA (BM-DA), which is similar to the M-DA but with early offers arriving on a predetermined day. We show that the BM-DA dominates the other two mechanisms in terms of student welfare. The reason is that early offers, by informing a student of the availability of certain programs, help her learn more efficiently. Moreover, we find that the student may do worse under the M-DA than under the DA because an early offer from an *ex ante* low-quality program may lead her to suboptimally accept it.

The comparison of the three mechanisms is further quantified in a set of simulations that use the same data from Germany. Relative to the DA, the BM-DA mechanism makes 72.2% of the students better off and only 20.8% worse off. The BM-DA also improves on the M-DA by making 41.7% of the students better off and 37.1% worse off.

Finally, our study reveals a previously unnoticed advantage of the program-proposing GS algorithm over the student-proposing version. Often, the latter is promoted because it is strategy-proof for students and achieves the student-optimal stable matching under the assumption that students have full information on their own preferences (see, e.g., Roth and Sotomayor 1990). The above dynamic multioffer mechanisms instead rely on the program-proposing GS to generate early offers, so that offers in the process cannot be withdrawn later. By contrast, the student-proposing GS is incompatible with a multioffer mechanism because it allows each program only to tentatively accept a student before the algorithm concludes and thus does not generate offers until the very end.

Other related literature.—Violations of the assumption that agents have full information on their own preferences have received little attention in the market design literature. As an exception, Narita (2018) finds that students reveal contradictory preferences in the main round and the subsequent round of school choice in New York City. Unlike our model, the author considers preference changes that are exogenous to the matching mechanism. Moreover, in the admissions to medicine programs in Germany, Dwenger, Kübler, and Weizsäcker (2018) document results implying a more complex process of preference formation than is commonly assumed.

Costly acquisition of information about preferences is a recent topic in the matching literature. For example, Chen and He (2021a, 2021b) theoretically and experimentally investigate students' incentives to acquire information about their own and others' preferences in school choice. These authors, as well as Artemov (2021), study the canonical DA and the Boston immediate-acceptance mechanisms. Similar to our paper, Immorlica et al. (2020) and Hakimov, Kübler, and Pan (2021) investigate how advising students on their admission chances can facilitate information acquisition, although these studies focus on single-offer mechanisms such as the iterative implementations of the DA and investigate the provision of historic information. Informing students of their admission chances can also help students beyond preference discovery, as students in centralized school choice may have incorrect beliefs about their admission chances and thus adopt suboptimal application strategies (He 2017; Kapor, Neilson, and Zimmerman 2020).

Sequential learning in our model is related to the theoretical literature on consumer search (Weitzman 1979; Ke, Shen, and Villas-Boas 2016; Doval 2018; Dzyabura and Hauser 2019; Ke and Villas-Boas 2019). In contrast to our setting, agents in that literature hold "offers" from all programs at the outset. The early-offer effect that we document resembles the increased demand for the item that is made salient in the model by Gossner, Steiner, and Stewart (2021), although an early offer in our model changes an agent's behavior by informing her that the corresponding program is available.

Our study complements the recent literature on dynamic implementations of single-offer mechanisms. Under the assumption of full information on own preferences, a series of papers—some of which are motivated by university admission procedures in practice (Bó and Hakimov 2018; Gong and Liang 2020)—investigate dynamic (or iterative) versions of the DA mechanism (Echenique, Wilson, and Yariv 2016; Klijn, Pais, and Vorsatz 2019; Bó and Hakimov 2020).⁶ Inspired by a university admission procedure used in France from 2009 to 2017, Haeringer and Iehlé (2021) study multiperiod admissions where a matching is computed in each period and students have the option to either finalize their matches or participate in the next period.

Our paper is also related to the literature on mechanism design with transferable utility where, in an initial stage, agents acquire information (e.g., Bergemann and Välimäki 2002), make investments (e.g., Hatfield

⁶ One of the main findings in this literature is that agents are more likely to report their true preferences under a dynamic DA than under a static one. There is a growing literature showing that under the static DA, many agents do not report their true preferences in the laboratory (Chen and Sönmez 2006; Rees-Jones and Skowronek 2018; for a survey, see Hakimov and Kübler 2021) and in the field (Artemov, Che, and He 2017; Shorrer and Sóvágó 2017; Hassidim, Romm, and Shorrer 2021).

et al. 2014; Nöldeke and Samuelson 2015), or face a surplus sharing rule (e.g., Dizdar and Moldovanu 2016). The sequential information updates owing to early offers in our setup resemble those in open versus closed auctions (see, e.g., Compte and Jehiel 2007).

The evolution of preferences in a decision process has been studied outside the market design literature. For example, Elster (1983) considers agents adjusting their preferences according to what is available to them, for which Alladi (2018) provides experimental evidence.

There is a large empirical literature on the determinants of college choice (for an early contribution, see, e.g., Manski and Wise 1983), which investigates the determinants of preferences rather than the process of preference discovery. Early admission offers play an important role in college admissions in the United States (Avery, Fairbanks, and Zeckhauser 2003; Avery and Levin 2010). Colleges seek to admit students who are enthusiastic about attending, and early admission systems give students an opportunity to signal their enthusiasm.

Organization of the paper.—After introducing the institutional background of Germany's university admissions in section I, we proceed to the data analysis and present the results as well as robustness checks in section II. To explain the findings, section III discusses several hypotheses and shows evidence from a survey. Section IV develops a model of university admissions with preference discovery that can explain the observed early-offer effect. We present the BM-DA mechanism and compare the DA, the M-DA, and the BM-DA in a theoretical model and simulations. Finally, section V concludes.

I. Institutional Background

A. University Admissions in Germany

Access to higher education in Germany is based on the principle that every student who completes the school track leading to the university entrance qualification (*Abitur*) should be given the opportunity to study in a university program of her choice. However, starting in the 1960s, a steep increase in the number of applicants created an overdemand for seats in programs such as medicine, so selection based on the final grade in the *Abitur* (numerus clausus) was introduced. In response to court cases brought forward against the universities, a central clearinghouse, the Zentralstelle für die Vergabe von Studienplätzen (ZVS), was established in 1972 to guarantee "orderly procedures."

In the 1990s and early 2000s, the number of programs administered through the ZVS clearinghouse steadily declined. The main reasons were that universities wanted to gain control of their admission process and that new bachelor's programs were created as part of the Bologna reforms

in fields of study that were not required for participation in the ZVS. By 2005, the only programs administered by the ZVS were in the fields of medicine, pharmacy, dental medicine, veterinary medicine, and psychology (the latter only until 2010–11). Seats for these programs were allocated according to a procedure involving quotas regulated by law.⁷ At the same time, severe congestion for many other programs emerged.

A reorganization and renaming of the clearinghouse from ZVS to Stiftung für Hochschulzulassung (literally, Foundation for University Admission) was completed in 2008, and the DoSV, a new admission procedure for programs other than medicine and medicine-related subjects, was implemented in 2012. Universities have the option to participate in the DoSV, and they can do so for a subset of their programs. Since 2012, the number of programs participating in this procedure has steadily increased.⁸

B. The DoSV: A Multioffer Procedure

The DoSV procedure has multiple phases. The early phases extend over several weeks and allow for student-program interactions that resemble those in a decentralized market. Each student applies to up to 12 programs and forms an ROL, without the possibility of later adding other programs to this list. Importantly, a student is not required to finalize her ROL until a later date, by which she may have received some offers from the programs to which she applied. With the finalized ROLs, the program-proposing GS algorithm is run to determine the matching. This procedure provides us with a unique data set of offers to students, their reranking of programs, and offer acceptances and rejections in real time. It allows us to measure the effects of early offers on the final admission outcome.

Figure 1 describes the timeline of the DoSV. The dates indicated are relevant for the winter term and are the same every year. We use data from the winter term, since admission for the summer term is possible only for a small number of programs. The procedure consists of the following phases.

Preparation phase (March 15–April 14).—The participating university programs register with the clearinghouse.

Application phase (April 15–July 15).—Students apply to at most 12 programs by directly submitting their application to the universities. The universities transmit to the clearinghouse the applications they have received

⁷ For analyses of the ZVS procedure, see Braun, Dwenger, and Kübler (2010), Westkamp (2013), and Braun et al. (2014).

⁸ For the winter term of 2015–16, 89 universities with 465 programs participated, compared with 17 universities with 22 programs in the first year in which the DoSV was implemented (winter term of 2012–13). The total number of students who were assigned to a program through the DoSV in 2015–16 was 80,905, relative to a total of 432,000 students who started university in Germany that year.




FIG. 1.—Timeline of the DoSV procedure (winter term). Our main interest is in phase 1—consisting of coordination phase 1 and the decision phase—where early offers are made and in phase 2—consisting of coordination phase 2—where the GS algorithm is run.

for their programs. By default, the programs in a student's ROL are ordered according to the arrival time of the corresponding applications at the clearinghouse. Students can, however, actively modify their ROL at any point during this phase. No new applications can be submitted after this phase, and no program will consider a student who did not apply to it.

Coordination phase 1 (July 16–August 15).—The programs rank applicants following a quota system. The size, number, and nature of the quotas are determined by state laws and regulations and by the universities themselves.⁹ For each program, the university admissions office creates rankings over its applicants—one for each quota—and transmits them to the clearinghouse. Every applicant is considered and ranked for every quota.¹⁰ Given the legal restrictions, universities have only very limited scope for strategic manipulation of their rankings of applicants,¹¹ although they may have some control as to when the lists are transmitted. Via automated emails, the clearinghouse sends admission offers for a program to its top-ranked students up to the program's capacity under each quota. We define these offers as early offers. Note that an early offer cannot be rescinded and does not expire until August 29 (the end of coordination phase 2). A student with one or more offers may accept one of them and leave the procedure, or she can choose to hold on to these offers (either all of them or a

⁹ The main quotas are (i) the *Abitur* quota (*Abiturbestenquote*), where the ranking is based on a student's average *Abitur* grade; (ii) the waiting time quota (*Wartezeitquote*), which gives priority to applicants who have waited for the greatest number of semesters since obtaining the *Abitur*; and (iii) the university selection quota (*Auswahlverfahren der Hochschulen*), which employs criteria that are determined by the universities themselves. Details on the process through which programs rank applicants and the role of quotas are provided in app. A.1 (apps. A–E are available online). These quotas are filled simultaneously, unlike in the old procedure for medical programs (see n. 3).

¹⁰ For this reason, we sometimes consider a student's most favorable quota at a program under which she has the highest probability of receiving an offer. A student cannot receive two offers from a program even if they are from different quotas.

¹¹ The *Abitur* and waiting time quotas rely on nonmanipulable criteria, while the ranking under the university selection quota is in practice almost entirely determined by students' *Abitur* grade. In our data, the average correlation coefficient between the rankings under the *Abitur* quota and the university selection quota is 0.86 across programs.

subset). If an offer is rejected, a new offer to the next applicant in that program's ranking is automatically generated. Students are informed about how a program ranks them in each quota, the total number of seats available in the program, and the number of students ranked above them in each quota who are no longer competing for a seat.

Decision phase (August 16–18).—Starting on August 16, university programs can no longer submit their rankings of applicants to the clearinghouse. However, early offers continue to be generated until August 18 because students may still reject offers received in coordination phase 1. Students are encouraged to finalize their ROL.

Coordination phase 2 (August 19–29).—At the beginning of this phase, a program may have (a) seats taken by students who have accepted its early offer and left the procedure, (b) seats/offers tentatively held by students who have kept their early offer from this program but have chosen to stay on, and (c) seats available because of its rejected early offers. Meanwhile, a student may have (i) left the procedure by accepting an early offer, (ii) kept an early offer and chosen to stay on with her final ROL, (iii) received no offer and stayed on with her final ROL, or (iv) exited the procedure, thereby rejecting all offers. Taking the remaining students and the seats available or tentatively held, the clearinghouse runs a program-proposing GS algorithm as follows:

1. Following the ranking of its applicants provided in coordination phase 1, a program sends admission offers to students ranked at the top of its ranking up to the number of available seats that are not tentatively held. However, the students who previously received an offer from the program can never receive the same offer again.
2. Students with multiple offers keep the highest-ranked one according to their final ROLs and reject all other offers. All other students are inactive.
3. Steps 1 and 2 are repeated until every program either has no seats left or has no more students to make offers to. Then each student is assigned to the program she holds, if any.

This phase uses the program-proposing GS algorithm where students start out with their highest-ranked offer from previous phases. They can never do worse than this offer, while a better offer may arrive. We define the GS algorithm as a computer code that uses information on students and programs to calculate a matching outcome. Importantly, the GS algorithm is silent about how students and programs provide such information. See appendix D.1 for more details and the definition of the GS algorithm.

Clearing phases 1 and 2 (August 30–September 4; September 30–October 5).—A random serial dictatorship is run to allocate the remaining seats to students who have not yet been admitted.

For the analysis of students' choices, we focus on coordination phase 1 and the decision phase. Since the two phases do not differ from the students' perspective, we group them together and call them phase 1 (see fig. 1). Coordination phase 2 is also of interest, and we call it phase 2. Recall that offers in phase 1 are early offers. We define *initial ROL* as the ROL over programs that the clearinghouse has recorded for each student at the beginning of phase 1, while the ROL at the end of phase 1 is defined as the *final ROL*. A program is defined as feasible to a student if the student applied to the program and was ranked higher than the lowest-ranked student who received an offer from the program in coordination phase 2.¹² A student may not actually receive an offer from a feasible program, as she might have left the procedure before she could receive the offer.

In contrast to single-offer mechanisms that typically deliver at most one offer to a student, the DoSV is a dynamic multioffer mechanism because a student may hold multiple offers in phase 1. Hence, we call this phase the multioffer phase.

C. Data

Our data set covers the DoSV procedure for the winter term of 2015–16. There are 183,088 students who applied to 465 programs at 89 universities in total. The data contain students' sociodemographic information (gender, age, postal code) and their *Abitur* grade.¹³ Furthermore, we observe students' ROLs at any point in time, the programs' rankings of applicants, the offers made by the programs throughout the procedure, the acceptance and rejection of offers by students, and the final admission outcome.

We exclude students with missing sociodemographic information as well as those who apply to specific programs with complex ranking rules. These are mostly students who want to become teachers and have to choose multiple subjects (e.g., math and English). For our analysis, we focus on the subsample of students who apply to at least two feasible programs and accept an offer. This leaves us with 21,771 students.

Table 1 provides summary statistics. On average, applicants to standard programs apply to 2.9 programs (col. 1). The corresponding figure is 4.2 among the 59% of students who apply to more than one program (col. 2)

¹² Our definition of feasibility is conditional on a student applying to a program. By contrast, the definition of feasibility in other papers on school choice and university admissions (e.g., Fack, Grenet, and He 2019) extends to the programs that a student did not apply to. This alternative is less appropriate in our setting because not all programs participate in the clearinghouse and our analysis is conditional on the programs to which each student has applied.

¹³ In the data, the *Abitur* grade is missing for about half of the students, but we can infer it for approximately two-thirds of applicants with a missing grade based on how students are ranked under the programs' *Abitur* quota. See app. A.1 for details.

TABLE 1
SUMMARY STATISTICS OF DoSV APPLICATION DATA FOR 2015–16 (WINTER TERM)

	SAMPLE		
	All Applicants to Standard Programs (1)	Applied to >1 Program (2)	Applied to ≥ 2 Feasible Programs and Accepted an Offer (3)
A. Students:			
Female	.579	.596	.558
Age	20.8 (3.2)	20.5 (2.6)	20.7 (3.1)
<i>Abitur</i> percentile rank (between zero and one)	.50 (.29)	.51 (.29)	.65 (.28)
B. Applications:			
Length of initial ROL (on July 15)	2.9 (2.6)	4.2 (2.7)	4.7 (2.9)
Reranked programs before phase 1 ^a	.537	.209	.998
Reranked programs during phase 1 ^b	.178	.305	.419
Fraction of programs located in student's municipality	.205 (.379)	.153 (.311)	.184 (.342)
Fraction of programs located in student's region (Land)	.623 (.446)	.610 (.420)	.583 (.417)
Average distance to ranked programs (km)	111 (127)	120 (119)	126 (122)
Top-ranked program (on July 15)—field of study: ^c			
Economics and business administration	.368	.397	.427
Psychology	.197	.204	.138
Social work	.121	.110	.044
Law	.110	.125	.170

Math/engineering/computer science	.065	.052	.097
Natural sciences	.055	.046	.059
Other	.085	.065	.066
C. Feasible programs and offers received:			
At least one feasible program	.505	.581	1.000
Received one or more early offers in phase 1	.475	.549	.989
D. Admission outcome:			
Canceled application before phase 2	.054	.042	.000
Accepted an early offer in phase 1	.220	.247	.554
Not initially top ranked	.262	.399	.369
Participated in phase 2	.725	.711	.446
Accepted an offer in phase 1 or phase 2	.448	.518	1.000
Number of days between offer arrival and acceptance	9.19	9.62	9.11
	(8.73)	(8.75)	(8.30)
Number of students	110,781	64,876	21,711

NOTE.—Summary statistics are computed from the DoSV data for the winter term of 2015–16. The main sample (col. 1) is restricted to students with nonmissing values and excludes those who applied to specific multicourse programs (*Mehrfachstudiengang*), which consist of two or more subprograms with complex assignment rules. Column 2 further restricts the sample to students who applied to two programs or more. Column 3 considers students who applied to at least two feasible programs and either actively accepted an early offer during phase 1 or were assigned to a program through the computerized algorithm in phase 2. For “*Abitur* percentile rank,” a higher value indicates a better *Abitur* grade. The distance between a student’s home and a program is computed as the Cartesian distance between the centroid of the student’s postal code and the geographic coordinates of the university in which the program is located. The number of days elapsed between offer arrival and acceptance is the number of days between the date the offer that was ultimately accepted was made to a student and the date it was accepted; for students who were assigned to their best offer by the computerized algorithm in phase 2, the acceptance date is set to the first day of phase 2—i.e., August 19, 2015. Standard deviations are shown in parentheses.

^a A student is considered as having reranked programs before phase 1 if she applied to only one program or if she manually altered the ordering of her applications before July 15.

^b A student is considered as having reranked programs during phase 1 if either the final ROL is different from the initial ROL or the student accepted an early offer from a program that she did not initially rank in first position.

^c For programs combining multiple fields of study, each field is assigned a weight of $1/k$, where k represents the number of fields.

and 4.7 among those who apply to at least two feasible programs and accept an offer (col. 3). Panel C reveals that 58.1% of students who apply to at least two programs (col. 2) have at least one feasible program; that is, they would have received at least one offer in the course of the procedure if they had not exited before phase 2. Importantly for our analysis, more than half of the students who apply to more than one program receive one or more offers in phase 1, and around one-quarter (24.7%) accept an offer in phase 1 (panel D). Among them, almost 40% accept an offer that was not their first choice in their initial ROL. Table 1 also indicates that only half of the students end up accepting an offer from a program in either phase 1 or phase 2.¹⁴ The rest may have accepted offers from programs that did not participate in the DoSV procedure.

D. Timing of Activities in the DoSV 2015–16

Figure 2 presents an overview of the activities in the DoSV procedure for 2015–16. It displays the points in time when students register with the clearinghouse, when they submit an ROL that is not changed any more (“finalize their ROL”), when they receive their first offers from programs (“receive an offer”), and when they exit the procedure. An important takeaway is that the first offers received by students are spread out over phase 1 (see also fig. B1; figs. B1–B3, C1, and C2 are available online). It is exactly this dispersed arrival of offers that allows us to identify the effect of early offers on offer acceptance. We show below that the offer arrival is not correlated with students’ initial ROLs and that early offers are not, on average, made by more selective or desirable programs. Instead, the time at which programs submit their rankings to the clearinghouse is determined by administrative processes within the universities.¹⁵

Almost all student exits from the DoSV take place in phases 1 and 2. During phase 1, students leave when they either accept an offer or cancel all applications. The number of exits peaks at the beginning of phase 2 when the clearinghouse automatically accepts an early offer from the top-ranked program of students who have not actively accepted the offer. The second spike occurs at the end of this phase, indicating that around half of the students do not receive any offer and therefore remain in the procedure until the very end.

¹⁴ Throughout the analysis, a student is defined as having accepted an offer when she either actively accepts an early offer in phase 1 or is assigned by the computerized algorithm in phase 2.

¹⁵ According to the Stiftung für Hochschulzulassung, the time point when the rankings are transmitted to the clearinghouse depends on the number of personnel available in an admissions office, the number of programs a university administers through the DoSV, the number of incomplete applications received, internal processes to determine the amount of overbooking for each program, and a university’s policy as to whether to check all applications for completeness or instead to accept applicants conditionally.

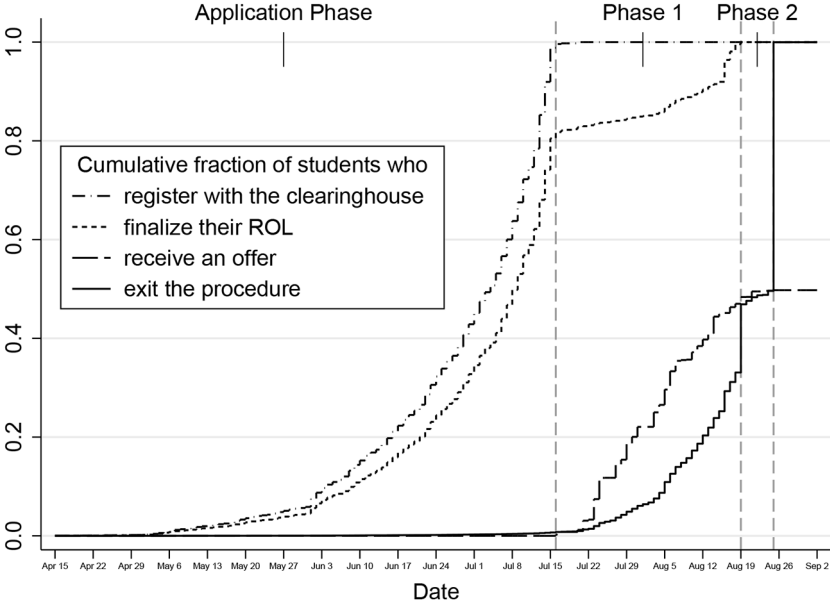


FIG. 2.—Activities during the DoSV procedure (winter term of 2015–16). This figure displays the evolution of several key indicators throughout the DoSV procedure for the winter term of 2015–16: (i) cumulative fraction of students who register with the clearinghouse during each phase (dash-dotted line), (ii) cumulative fraction of students who finalize their ROL of programs (short-dashed line), (iii) cumulative fraction of students who receive at least one offer (long-dashed line), and (iv) cumulative fraction of students who exit the procedure (solid line) due to one of the following motives: active acceptance of an early offer during phase 1, automatic acceptance of the best offer during phase 2, cancellation of application, rejection due to application errors, or rejection in the final stage for students who participate in phase 2 but receive no offer.

Next, we disaggregate the exits by their reason for leaving the procedure. Figure 3 shows that 22% of students actively accept an offer during phase 1, 22% receive their best offer during phase 2 when the GS algorithm is run (of which two-thirds are automatically removed on the first day because they have an offer from their top-ranked program), and 14% cancel their applications at some point, while the remaining 41% participate in phase 2 but receive no offer.

II. Accepting Early Offers: Empirical Results

We now turn to the question of whether a program’s offer to a student is more likely to be accepted when it is an early offer. Our analysis is restricted to a student’s feasible programs. Recall that a program is feasible to a student if the student applied to the program and would have received its offer, provided that she remained in the procedure until phase 2 while




FIG. 3.—Reasons for exiting the DoSV procedure (winter term of 2015–16). This figure shows the cumulative admission outcomes of students throughout the DoSV procedure for the winter term of 2015–16: (i) cumulative fraction of students who actively accept an early offer received during phase 1 (area with horizontal hatching), (ii) cumulative fraction of students on whose behalf the clearinghouse accepts their best offer during phase 2 (area with diagonal hatching), (iii) cumulative fraction of students who cancel all applications (dotted area), and (iv) cumulative fraction of students who participate in phase 2 but receive no offer (area with vertical hatching).

holding other students' behavior constant. Infeasible programs are irrelevant to a student's offer acceptance decision. If offers from the feasible programs arrive in an exogenous sequence, common matching models predict no early-offer effect on acceptance because of the assumption that students have full information on their own preferences.

A. Empirical Approach

We adopt a logit model to assess whether students are more likely to accept early offers. Let F_i represent student i 's set of feasible programs, with the programs being indexed by j . For j in F_i ,

$$U_{i,j} = \mathbf{Z}_{i,j}\beta + \epsilon_{i,j},$$

where $U_{i,j}$ represents how student i values feasible program j at the time of making the ranking or acceptance decision, $\mathbf{Z}_{i,j}$ is a row vector of student-program-specific characteristics, and $\epsilon_{i,j}$ is independent and identically

distributed (i.i.d.) according to a type I extreme value (Gumbel) distribution. Note that U_{ij} may not be i 's true, full-information utility but rather stands for i 's valuation of j conditional on all her information at the time of the decision.

Students accept the offer from their most preferred feasible program. Therefore, we can write i 's choice probability for $j \in F_i$ as

$$\begin{aligned} & \mathbb{P}\left(i \text{ accepts program } j\text{'s offer} | F_i, \{\mathbf{Z}_{ij}\}_{j \in F_i}\right) \\ &= \mathbb{P}\left(U_{ij} \geq U_{ij'}, \forall j' \in F_i | F_i, \{\mathbf{Z}_{ij}\}_{j \in F_i}\right) = \frac{\exp(\mathbf{Z}_{ij}\beta)}{\sum_{j' \in F_i} \exp(\mathbf{Z}_{ij'}\beta)}. \end{aligned} \quad (1)$$

In effect, our analysis assumes that a student ranks her most preferred feasible program above all other feasible programs in her final ROL, conditional on the set of applications that were already finalized in the application phase.¹⁶ By focusing on feasible programs, we allow for the possibility that a student arbitrarily ranks an infeasible program without any payoff consequences (Artemov, Che, and He 2017; Fack, Grenet, and He 2019).

To investigate whether receiving an early offer from program j in phase 1 (as opposed to receiving it in phase 2) increases the probability that i accepts j 's offer, the model is specified as

$$U_{ij} = \theta_j + \delta \text{ Early Offer}_{ij} + \gamma d_{ij} + \mathbf{X}_{ij}\lambda + \epsilon_{ij}, \quad (2)$$

where θ_j represents the fixed effect of program j and d_{ij} represents the distance between student i 's postal code and the address of the university where the program is located. EarlyOffer_{ij} is a dummy variable that equals one if i receives a potential early offer from program j during phase 1 (i.e., up to August 18) rather than in phase 2. A potential early offer is defined as either an offer that was actually received by the student in phase 1 or—if the student canceled her application to the program before phase 2—an offer that the student would have received in phase 1.¹⁷ The row vector \mathbf{X}_{ij} includes other student-program-specific controls. The coefficient of interest, δ , is thus identified by the variation in the timing of the offer arrival across students and programs, conditional on the programs' observed heterogeneity and unobserved average quality.

¹⁶ As the DoSV is based on the program-proposing GS algorithm, it is not strategy-proof for students. A student can misreport her preferences in an ROL by reversal (i.e., reversing the true preference order of two programs) and dropping (i.e., dropping programs from the true preference order). To identify profitable misreports, a student usually needs rich information on other students' and programs' preferences. In a low-information environment, Roth and Rothblum (1999) show that reversals are not profitable. Dropping does not create an issue for our analysis because we focus on the programs ranked in an ROL.

¹⁷ Our focus on potential rather than actual early offers ensures that our definition of an early offer does not depend on the student's acceptance decision. In the data, 90% of potential early offers were actually received by students.

B. Identifying Assumption: Exogenous Arrival of Offers

One potential concern is that early offers might be more attractive than those arriving later for reasons unrelated to their arrival time. The specification in equation (2) requires that $\text{EarlyOffer}_{i,j}$ is independent of shocks ($\epsilon_{i,j}$) conditional on other controls. To test this identifying assumption, we examine whether the average quality of potential offers varies over time using two measures of a program's selectivity and desirability. After calculating these measures for every program, we take the average over all offers that are sent out on a given day (weighted by the number of offers made by each program).

In figure 4A, the selectivity measure is the *Abitur* percentile (between zero and one, with a higher value indicating a higher selectivity) of the last student admitted to the program in phase 2, with the percentile computed among all DoSV applicants. A higher value on this measure indicates a higher degree of selectivity. In figure 4B, similar to Avery et al. (2013), a program's desirability is inferred from students' acceptance decisions among their feasible programs. Specifically, we estimate the conditional logit model described in equation (2) and use the program fixed effect estimates (standardized to have a mean of zero and a unit variance across programs) as a proxy for program desirability.¹⁸

In either of the two figure panels, there is no clear pattern over time, which is consistent with the timing of offers being determined mostly by the universities' administrative processes rather than by strategic considerations. If anything, the very early offers tend to come from slightly less selective or less desirable programs.

We further investigate time trends in offers based on regression analyses. Column 1 of table 2 (panel A) shows an insignificant correlation between the selectivity measure on each day and the number of days that have elapsed since the start of phase 1. Column 3 repeats the same regression for the desirability measure. The correlation is positive and significant, implying that earlier offers are from marginally less desirable programs. Columns 2 and 4 regress the selectivity measures on week dummies. Indeed, the results show that the very early offers are less desirable. Panel B further investigates time trends in the variance of the selectivity/desirability of offers, since the higher acceptance rate of early offers may be caused by their more dispersed quality relative to later offers (even though the

¹⁸ We use the results in col. 3 of table 3, in which we control for a quadratic function of distance to the program, whether the program is in the student's region (Land), and how the program ranks the student. We also considered an alternative measure of program desirability based on students' initial ROLs rather than on their acceptance decisions, using a larger sample of students who applied to at least two programs (not necessarily feasible). The results are very similar to those based on students' acceptance decisions. Although positive and statistically significant, the correlation between the two measures of program selectivity and desirability is small (0.17), suggesting that these proxies capture different dimensions.

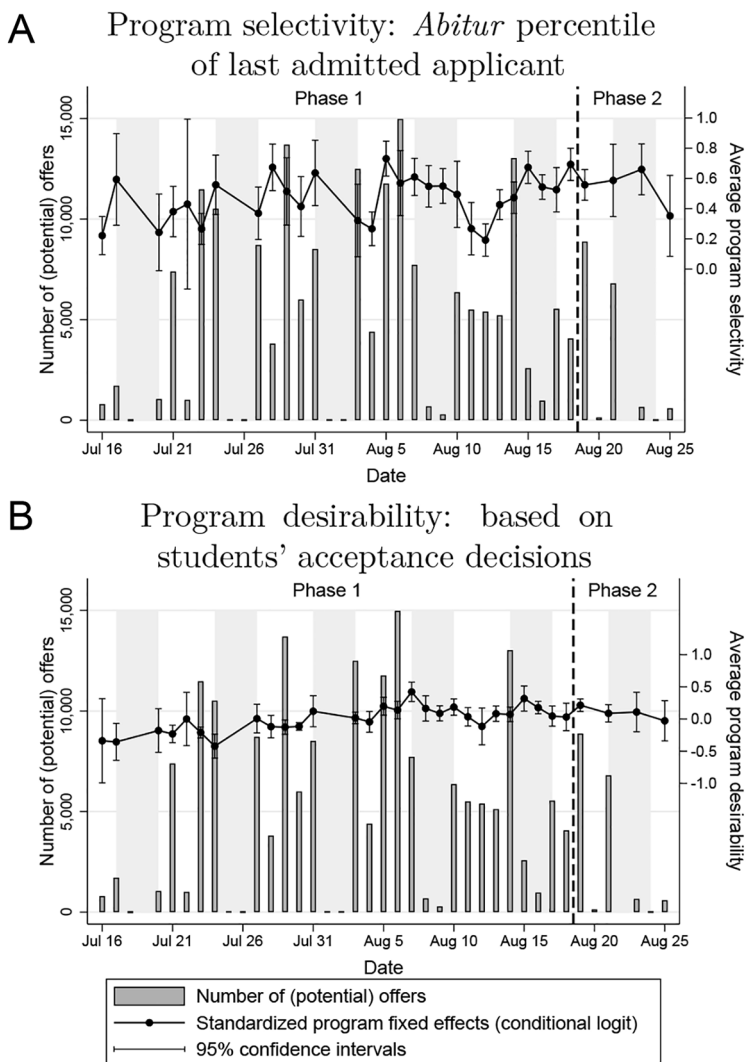


FIG. 4.—Offer arrival, program selectivity, and program desirability. Vertical bars indicate the number of potential offers sent out by programs on a given day throughout the DoSV procedure (winter term of 2015–16). Potential offers are defined as either actual offers that were sent out to students or offers that a student would have received had she not canceled her application to the program. The jagged line shows the average selectivity/desirability of programs sending out offers on a given day, with 95% confidence intervals denoted by vertical T bars. In A, a program's selectivity is proxied by the *Abitur* percentile (between zero and one, with a higher value indicating a higher selectivity) of the last student admitted to the program in phase 2. In B, a program's desirability is proxied by the program fixed effect estimates (standardized to have a mean of zero and a unit variance across programs) from the conditional logit model whose results are shown in column 3 of table 3. The selectivity/desirability of potential offers made on a given day is computed as the average over the programs making these offers on a given day (weighted by the number of offers made by each program). The measures are not shown for days in which fewer than 150 potential offers were made, which mostly coincide with weekends (denoted by gray-shaded areas).

TABLE 2
OFFER ARRIVAL, PROGRAM SELECTIVITY, AND PROGRAM DESIRABILITY: REGRESSION ANALYSES

	PROGRAM SELECTIVITY: <i>ABITUR</i> PERCENTILE OF LAST ADMITTED APPLICANT		PROGRAM DESIRABILITY: BASED ON STUDENTS' ACCEPTANCE DECISIONS	
	(1)	(2)	(3)	(4)
A. Dependent Variable: Average Selectivity/Desirability of Programs Making Offers on a Given Day				
Days since start of phase 1	.0038 (.0024)		.0133*** (.0028)	
Week of (potential) offer arrival:				
Week 1 (July 16–22)		–.110 (.104)		–.294*** (.061)
Week 2 (July 23–29)		–.052 (.116)		–.248*** (.088)
Week 3 (July 30–August 5)		Reference week		Reference week
Week 4 (August 6–12)		–.028 (.119)		.090 (.088)
Week 5 (August 13–19)		.027 (.103)		.053 (.064)
Week 6 (August 20–25)		.066 (.101)		.011 (.056)
B. Dependent Variable: Variance of Selectivity/Desirability of Programs Making Offers on a Given Day				
Days since start of phase 1	.0008 (.0006)		.0029 (.0034)	
Week of (potential) offer arrival:				
Week 1 (July 16–22)		–.020 (.020)		–.198 (.119)
Week 2 (July 23–29)		–.008 (.017)		–.199 (.119)
Week 3 (July 30–August 5)		Reference week		Reference week
Week 4 (August 6–12)		.005 (.024)		–.008 (.197)
Week 5 (August 13–19)		–.005 (.018)		–.158 (.118)
Week 6 (August 20–25)		.067*** (.023)		–.199 (.118)
Days in phase 1	39	39	39	39
Number of potential offers	192,840	192,840	192,840	192,840

NOTE.—This table reports linear regression estimates to test whether the timing of offers is correlated with the selectivity (cols. 1, 2) or desirability (cols. 3, 4) of the programs sending out these offers. The sample includes all potential offers, i.e., offers that either were sent out to students or could have been sent out had the students not canceled their application to the program. The day of arrival of each (potential) offer is identified as the day it became feasible to the student. In cols. 1 and 2, a program's selectivity is proxied by the *Abitur* percentile (between zero and one, with a higher value indicating a higher selectivity) of the last student admitted to the program in phase 2. In cols. 3 and 4, a program's desirability is proxied by the program fixed effect estimates (standardized to have a mean of zero and a unit variance across programs) from the conditional logit model whose results are shown in col. 3 of table 3, using the sample of students who applied to at least two feasible programs and accepted an offer. After calculating the selectivity/desirability measures for every program, we compute the mean and variance of each measure over all offers that were sent out on a given day (weighted by the number of offers made by each program). In cols. 1 and 3, the program selectivity/desirability measures are regressed on a linear time trend; in cols. 2 and 4, they are regressed on a vector of week dummies, with the third week of the DoSV procedure (July 30–August 5) as the omitted category. Standard errors clustered at the day level are shown in parentheses.

*** $p < .01$.

average quality is the same). We find no evidence of such trends: the coefficients on days elapsed or week dummies are almost never statistically significant.

Additionally, in table 4 (cols. 1–3) we show that early offers are not correlated with how students rank feasible programs in their initial ROLs. Taken together, the results indicate that programs from which students receive early offers are not more attractive and that they were initially not ranked higher by students.

C. Empirical Results on the Early-Offer Effect

We use the sample of students who applied to at least two feasible programs and either actively accepted an early offer in phase 1 or were assigned by the GS algorithm in phase 2. In the empirical analysis, we refer to these students as having accepted a program's offer. In total, there are 21,711 such students. Together, they applied to 66,263 feasible programs.

We start with the specification in equation (2) to study the impact of early offers on the acceptance of offers. The regression results are reported in table 3. Column 1 includes the early-offer dummy (*EarlyOffer*) and program fixed effects. The program fixed effects capture observed and unobserved program-specific characteristics, such as selectivity or faculty quality, which might be correlated with students' offer acceptance decisions. The coefficient on *EarlyOffer* is positive and significant, suggesting that having received an early offer increases the probability of a student accepting that offer. Column 2 adds another dummy variable that is equal to one for the first early offer (*FirstEarlyOffer*).¹⁹ Students are even more likely to accept the first offer, while all early offers remain more likely to be accepted than other offers. The results are qualitatively similar when we add further controls, such as a quadratic function of distance to the program and whether the program is in the student's region (*Land*; col. 3), how the program ranks the student (col. 4),²⁰ and the chances of a student not receiving an offer from the program in phase 2 (col. 5). We proxy the last control variable by the ratio between a student's rank and the rank of the last student who received an offer from the program in phase 1. The variable is zero if a student has an early offer from the program. We thus allow for the possibility that a student accepts an early offer because she does not expect to receive other offers in phase 2.

¹⁹ In our sample, the average time between the first and second early offers is 4.89 days among the 17,351 students who received two or more early offers. When we consider the 19,582 students who received or could have received two or more early offers, the average time between the first and second (potential) early offers is 5.38 days.

²⁰ We compute how a student is ranked by a given program using the student's percentile (between zero and one, with a higher value indicating a better rank) among all applicants to the program under the *Abitur* quota.

TABLE 3
EARLY OFFER AND ACCEPTANCE AMONG FEASIBLE PROGRAMS: CONDITIONAL LOGIT

	(1)	(2)	(3)	(4)	(5)
A. Estimates:					
EarlyOffer: potential offer from program in phase 1	.484*** (.041)	.410*** (.043)	.411*** (.044)	.404*** (.044)	.424*** (.108)
FirstEarlyOffer: first offer in phase 1		.133*** (.022)	.147*** (.023)	.147*** (.023)	.147*** (.023)
Distance to university (thousands of km)			-9.36*** (.33)	-9.37*** (.33)	-9.37*** (.33)
Distance to university squared			12.52*** (.55)	12.54*** (.55)	12.54*** (.55)
Program in student's region (Land)			-.005 (.039)	-.006 (.039)	-.006 (.039)
Program's ranking of student (between zero and one)				.439* (.227)	.442* (.227)
Chances of not receiving an offer from program in phase 2					.016 (.076)
Program fixed effects (376 programs)	Yes	Yes	Yes	Yes	Yes
Number of students	21,711	21,711	21,711	21,711	21,711
Total number of feasible programs	66,263	66,263	66,263	66,263	66,263

B. Marginal effects on acceptance probability of feasible programs:				
Nonfirst early offer (percentage points)	10.4 (1.5)	8.7 (1.4)	8.4 (1.6)	8.3 (1.5)
First early offer (percentage points)		11.6 (1.7)	11.5 (2.0)	11.3 (2.0)
C. Marginal effects on utility (measured in distance):				
Nonfirst early offer (km)			-59	-58
First early offer (km)			-78	-77

NOTE.—This table reports estimates from a conditional logit model for the probability of accepting a program among feasible programs. Each student's choice set is restricted to the feasible programs that she included in her initial ROL, i.e., the programs from which she could have received an offer by the end of Phase 2. A program's ranking of the student is computed as the student's percentile (between zero and one, with a higher value indicating a better rank) among all applicants to the program under the *Abitur* quota. The chances of not receiving an offer from a program in phase 2 are proxied by the ratio between the student's rank and the rank of the last student who received an offer from the program in phase 1; the variable is zero if the student has an early offer from the program. For panel B, the baseline (no early offer) acceptance probability is 38.5%. For the marginal effect of a nonfirst early offer on offer acceptance probability, we measure the difference between the following two predictions on offer acceptance behavior: while keeping all other variables at their original values, we let (i) EarlyOffer = 1 and FirstEarlyOffer = 0 and (ii) EarlyOffer = FirstEarlyOffer = 0. The baseline probability is the average of the second prediction across students, while the reported marginal effect is the average of the difference between the two predictions across students. The marginal effect of the first early offer is calculated in a similar manner. For panel C, the average distance to ranked programs is 126 km. The marginal effect of a nonfirst early offer measured in distance is calculated as the reduction in distance from 126 km that is needed to equalize the effect on utility of switching EarlyOffer from one to zero. A similar calculation is performed for the marginal effect of the first early offer. Standard errors are shown in parentheses.

* $p < .1$.
 *** $p < .01$.

All these results show a positive early-offer effect on offer acceptance and a larger effect for the first early offer. To quantify the effects, we calculate the impact of an early offer on the probability of offer acceptance (table 3, panel B). On average, a feasible program is accepted by a student with a probability of 0.385. An early offer that is not the first one increases the acceptance probability by 8.7 percentage points, or 27.9%, based on the estimates in column 5.²¹ This is of the same order of magnitude as the estimates from other specifications (cols. 1–4). The first early offer has a larger marginal effect. It increases the acceptance probability by 11.8 percentage points, or 38.3%, based on the estimates in column 5.²²

Another way to evaluate the magnitude of the early-offer effect is to use distance as a numeraire (table 3, panel C). At the sample mean of distance (126 km, as shown in col. 3 of table 1), an early offer that is not the first offer gives the program a boost in utility equivalent to reducing the distance by 61 km, based on the results in column 5.²³ The first early offer has an even larger effect, amounting to a reduction of 79 km.

We explore the heterogeneity of the early-offer effect by adding interactions between student characteristics and the *EarlyOffer* and *FirstEarlyOffer* dummies. The results are shown in table B1 (tables B1–B10 and C1 are available online). They indicate that female students respond less to early offers, although there is no additional heterogeneity in the effect of the first early offer. Students with a better *Abitur* grade are less responsive to the first early offer but do not behave differently for other early offers. The number of feasible programs that a student has does not change the early-offer and first-early-offer effects.

1. Early-Offer Effect on Students' Reranking Behavior

We extend the above analyses to students' reranking behavior, which is possible because our data record all changes in a student's ROL. Specifically, we use the initial and final ROLs of each student to investigate whether the ranking of programs is influenced by early offers.

²¹ This marginal effect is the difference between the following two predictions on offer acceptance: while keeping all other variables at their original values, we let (i) *EarlyOffer* = 1 and *FirstEarlyOffer* = 0 and (ii) *EarlyOffer* = *FirstEarlyOffer* = 0. The baseline probability is the average of the second prediction across students, while the reported marginal effect is the average difference between the two predictions across students. The marginal effect of the first early offer is calculated in a similar manner.

²² When we further include the dummies for the second and third early offers, we find that their marginal effects decrease monotonically in the arrival order.

²³ To calculate the marginal effect of this nonfirst early offer, we calculate the reduction in distance from 126 km that is needed to equalize the effect on utility of switching *EarlyOffer* from one to zero. A similar calculation is performed for the marginal effect of the first early offer.

As justified in section II.A, we restrict our attention to feasible programs and extract information from a final ROL as follows: (i) for a student who actively accepted an early offer during phase 1, we code only that she prefers the accepted offer to all other feasible programs in her ROL (clearly, we do not have credible information on the relative rank order among all the feasible programs); (ii) for a student who was assigned to a program in phase 2, we use the rank order among the feasible programs in her final ROL up to the first program that made her an early offer in phase 1. Programs ranked below this highest-ranked early offer are coded only to be less preferred than those ranked above. Their relative rank order is ignored because it is payoff-irrelevant for the student.

Using a rank-ordered logit (or exploded logit), we obtain the results in panel A of table 4. Columns 1–3 are placebo tests in which we use a student's initial ROL as the outcome variable. The results show that there is no significant correlation between the initial rank order of a program and receiving an early offer from that program, which is consistent with the finding that early offers are not from more attractive programs.

Columns 4–6 reveal that receiving an early offer induces a student to rank that program higher in her final ROL. Similarly, the first early offer enjoys a premium.

Using the estimation results, we further quantify the effect of receiving an early offer and find results that are almost identical to those for acceptance probability (table 3). Panel B of table 4 presents the marginal early-offer effects on the probability of top ranking the program in one's ROL among feasible programs. A nonfirst early offer increases the probability of top ranking the program by 8–10 percentage points, or 25%–30% (cols. 4–6).²⁴ The first early offer has a larger effect. It increases the probability of top ranking the program by 11 percentage points, or 33%–36% (cols. 5, 6).

In panel C of table 4, at the sample mean of distance (126 km), the first early offer gives the program a boost in utility that is equivalent to reducing the distance by 76 km, while the distance-equivalent utility of other early offers is a reduction of 60 km on average.

2. Robustness Checks

Our results are robust to a number of sensitivity tests. First, we consider a more restrictive definition of program feasibility to account for the possibility that students who barely cleared a program's admission cutoff in phase 2 may have considered this program infeasible *ex ante*. We relabel as infeasible any program j in student i 's initial ROL such that the ratio r_{ij} between the student's rank under the program's most favorable quota

²⁴ The calculations are similar to those on offer acceptance probability. See the details in n. 22.

TABLE 4
INITIAL VS. FINAL RANKING OF FEASIBLE PROGRAMS: RANK-ORDERED LOGIT

	INITIAL ROL (at Start of Phase 1)			FINAL ROL (at End of Phase 1)		
	(1)	(2)	(3)	(4)	(5)	(6)
A. Estimates:						
EarlyOffer: potential offer from program in phase 1	-.033 (.028)	-.028 (.028)	-.071 (.078)	.453*** (.040)	.387*** (.042)	.405*** (.105)
FirstEarlyOffer: First offer in phase 1		-.012 (.016)	-.003 (.016)		.118*** (.022)	.131*** (.023)
Distance to university (thousands of km)			-5.44*** (.21)			-9.15*** (.32)
Distance to university squared			7.21*** (.36)			12.17*** (.53)
Program is in student's region (Land)			.004 (.026)		.002 (.038)	
Program's ranking of student (between zero and one)			.130 (.155)		.448*** (.224)	
Chances of not receiving an offer from program in phase 2			-.018 (.056)			.019 (.074)
Program fixed effects (376 programs)	Yes	Yes	Yes	Yes	Yes	Yes

Number of students	21,711	21,711	21,711	21,711	21,711	21,711
Total number of feasible programs	66,263	66,263	66,263	66,263	66,263	66,263
B. Marginal effects on probability of ranking feasible program as top choice:						
Nonfirst early offer (percentage points)		9.7 (1.5)				
First early offer (percentage points)		8.3 (1.3)				
		10.8 (1.6)				
C. Marginal effects on utility (measured in distance):						
Nonfirst early offer (km)						-60
First early offer (km)						-76

NOTE.—This table reports estimates from a rank-ordered logit model for the probability of observing a student's initial and final ROLs of feasible programs. The sample and variables are the same as in table 3. Columns 1–3 consider students' initial ROLs, while cols. 4–6 consider their final ROLs. We take as a student's initial ROL the partial order of feasible programs that she ranked at the beginning of phase 1. In the analysis of final ROLs, we use only the payoff-relevant information contained in each ROL. For panel B, the baseline (no early offer) acceptance probability is 38.5%. For the marginal effect of a nonfirst early offer on the probability of top ranking a program, we measure the difference between the following two predictions on top-ranking behavior: while keeping all other variables at their original values, we let (i) EarlyOffer = 1 and FirstEarlyOffer = 0 and (ii) EarlyOffer = FirstEarlyOffer = 0. The baseline probability is the average of the second prediction across students, while the reported marginal effect is the average of the difference between the two predictions across students. The marginal effect of the first early offer is calculated in a similar manner. For panel C, the average distance to ranked programs is 126 km. This panel is calculated in the same way as in table 3. Standard errors are shown in parentheses.

** $p < .05$.

*** $p < .01$.

and the rank of the last student who received an offer from the program under that quota is between a prespecified value \bar{r} (<1) and 1.²⁵ The sample selection and the dummy for offer acceptance are adjusted according to the new feasibility. Table B2 summarizes the results from the conditional logit model when we vary the ratio \bar{r} between 0.5 and 1. The early-offer and first-early-offer effects are very similar to the baseline estimates across the values of \bar{r} . We also assess and confirm the robustness of our estimates to artificially expanding students' feasible sets (see table B3).

Second, we consider the possibly heterogeneous effect of early offers during the first 2 weeks of phase 1, since there is some evidence that the very early offers are from slightly less selective programs. Table B4 shows that our main results are not driven by these very early offers.

How a student ranks her feasible programs in the initial ROL may reflect her preferences. In the investigation of the early-offer effect on offer acceptance, we further control for how the student ranks each program. Table B5 reveals that this further control does not change our results.

In the analysis of student ranking behavior, one may be concerned that some students' initial ROLs may not be meaningful because students know that they have until the end of phase 1 to change them. In table B6, we restrict the sample to students with an initial ROL that they had reranked at some point in the application phase. For a student in this subsample, her initial ROL is more likely to reflect her initial preferences. The results in the table are very similar to those based on the full sample.

Finally, as additional evidence for our main findings, we implement an alternative empirical design to estimate the early-offer effect. We take advantage of the discontinuity around the cutoff rank that allows students to receive an offer in phase 1 (i.e., an early offer) rather than in phase 2 and compare the probability that a (potential) offer is accepted by students on either side of the cutoff. While intuitively appealing, this regression discontinuity (RD) design has several limitations in our setting, leading to our decision to not use it as our main empirical strategy.²⁶ Bearing in mind these limitations, we apply a fuzzy RD design to estimate the impact of receiving a (potential) early offer from a program on the acceptance

²⁵ Note that this ratio is the same as the one we use to proxy a student's chances of not receiving an offer from the program in phase 2 (table 3, col. 5).

²⁶ We identify three limitations of the RD design in our setting. First, in the context of the DoSV procedure, the RD design is fuzzy rather than sharp because students are ranked under multiple quotas for any given program (the average program has six quotas). Namely, a student who fails to clear the phase 1 cutoff under a program's quota q can receive an early offer from the same program under a different quota q' , thus reducing the discontinuity in the early-offer probability at the phase 1 cutoff under quota q . Second, it allows us to estimate only the early-offer effect on offer acceptance and not to compare the effects of the first versus subsequent early offers or to analyze students' reranking behavior. Third, it identifies the early-offer effect only for the subgroup of students who barely cleared or barely missed the phase 1 cutoff, whereas we are interested in estimating this effect for a broader population of applicants.

probability of students who were ranked just above versus just below the program's phase 1 cutoff rank (see app. C for details). Reassuringly, the results are very similar to those from the conditional logit model. The RD estimates indicate that receiving an early offer from a program increases the probability of accepting that program by 8–9 percentage points (see table C1).

III. Explanations of the Early-Offer Effect

In this section, we begin by testing and ruling out several possible explanations of the early-offer effect. Our preferred explanation is that students learn about their preferences over programs in the course of the procedure, which is corroborated in a student survey presented at the end of this section. Based on these findings, we develop a model with preference discovery in section IV that can explain the early-offer effect.

A. Alternative Explanations

Here we discuss alternative explanations of the early-offer effect.

Preference for being ranked high by a program.—A student may respond positively to an early offer if she thinks that a program is revealing its appreciation or a high match quality by making her an early offer.²⁷ This implies that students care about how a program ranks them. If the early-offer effect is driven by students' preferences for being ranked high by a program, the effect would disappear or decrease when we control for how a program ranks a student (which is observable to her in the DoSV). With this additional control, column 4 of table 3 shows that the early-offer effect remains the same.²⁸

Immediate reaction.—The early-offer effect can be driven by an immediate reaction—for example, from a feeling of relief.²⁹ However, we find some counterevidence. First, a nonfirst early offer is also more likely to be accepted than nonearly offers (table 3). Second, students do not immediately accept an offer upon its arrival. The average waiting time before accepting an offer is 9 days (fig. B2; table B7), and the distributions of offers and acceptances on each day of the week differ markedly (fig. B3),

²⁷ Relatedly, Antler (2019) studies a model in which workers experience a disutility when they are ranked low on the employer's preference list.

²⁸ Table B8 further shows that our results are robust to controlling for potential nonlinearities in the relationship between a student's rank and her acceptance probability (e.g., students dislike being at the bottom of their class). We perform this test either by including a second-/fourth-order polynomial of a student's rank (between zero and one) or by splitting each program's ranking of applicants into quartiles/deciles and including these indicators as controls. Our estimates remain essentially unchanged.

²⁹ Theories of dual selves posit that decisions are influenced by an intuitive system, system 1, that performs automated or emotion-driven choices and a deliberative system, system 2, responsible for more reflective decisions. See, e.g., Kahneman (2003, 2011), Loewenstein and O'Donoghue (2004), Fudenberg and Levine (2006), and Evans (2008).

with a significant fraction of acceptances occurring on the weekend when almost no offers are made.

Pessimism about future offers.—Students may accept early offers if they assign a close to zero probability of receiving a better offer later. However, given the information provided by the DoSV, it is likely that students have a reasonably accurate assessment of the offer probabilities. Taking a student's initial ROL as a proxy of her preferences, we find evidence that the probability of receiving a better offer in phase 2 is substantial. Among the students who have at least one (potential) offer in phase 2,³⁰ almost half rank a phase 2 offer above all early offers in their initial ROLs. These students could have received better offers in phase 2 had they remained in the system. Moreover, among all students in our sample, 65% have at least one potential offer in either phase that is ranked above their first early offer in their initial ROLs. Therefore, given the information available and the actual possibility of receiving better offers in the future, it seems unlikely that the early-offer effect would be driven by students' pessimistic beliefs.³¹

Head start in the housing market.—Student dormitories are scarce in Germany and can accommodate only a small fraction of students. Hence, a student may accept an offer early to have a head start when searching for housing. However, this incentive does not exist for students who attend a university in their own municipality and often live with their parents. We can therefore test this housing demand hypothesis by checking whether the early-offer effect is also observed in the subsample of students who applied only to local programs. The results are summarized in table B9. Despite the relatively small size of this subsample (11.3% of the full sample), the early-offer and first-early-offer effects are statistically significant and similar in magnitude to our baseline estimates.³² We thus conclude that housing concerns are unlikely to explain the early-offer effect.

Aversion to the computerized assignment.—Some students may dislike being assigned by a computerized algorithm in phase 2 and may therefore actively accept an offer in phase 1. We document two pieces of evidence to show that this is unlikely to drive our findings. First, almost half of the

³⁰ Those who do not have a (potential) phase 2 offer do not contribute to identifying the early-offer effect, although they do contribute to identifying the first-early-offer effect if they have multiple early offers.

³¹ Consistent with this interpretation, our earlier results show that the early-offer and first-early-offer effects are robust to adopting a more restrictive definition of program feasibility (see table B2, which excludes programs that students may not have considered as feasible).

³² As additional evidence, we show in table B10 that our results are robust to restricting the sample to students who did not accept an early offer until at least halfway through phase 1 (i.e., August 2) and hence who are less likely to have accepted an early offer in order to start looking for housing as soon as possible. The estimated early-offer effects are comparable to those using the full sample, while the first-early-offer effects are only slightly smaller when expressed in terms of marginal effect on acceptance probability.

students actively choose to participate in a computerized assignment in phase 2 by top ranking an early offer when entering phase 2 (see fig. 3). Second, the effect of the first early offer is larger than that of other early offers, which cannot be explained by an aversion to the computerized assignment.

Beyond the possible explanations discussed so far, we cannot rule out that early offers create an endowment effect that might cause students to value programs with early offers more than other programs (Kahneman, Knetsch, and Thaler 1990). Our model in section IV shows that such a behavioral factor is not necessary to explain the early-offer effect. Moreover, we have direct evidence for our model from a survey. It is presented below and documents that students learn about programs in the course of the DoSV procedure.

B. Evidence from a Survey

To obtain direct information on students' decision-making, we conducted a survey. It was administered by the Stiftung für Hochschulzulassung as part of an official survey that was accessible through a link on the website of the DoSV. Around 9,000 students completed it in 2015 (i.e., our sample period). Information about the setup of the survey and the complete list of questions are provided in appendix A.3.

Table 5 considers three different groups of survey respondents in panels A–C. Respondents in panel A are those who applied to more than one program, accounting for 66% of all respondents. Among them, 30% report that they did not have a clear ranking over programs at the time of application because they needed more research to form their preferences; 25% agree that coming up with a preference ranking was very difficult and that they wanted to delay the decision for as long as possible.

Among respondents who applied to more than one program, 84% had received at least one offer at the time of the survey, including offers from programs not in the DoSV. Panel B shows that receiving an offer increased the time spent on learning about that university for 61% of students, while only 10% spent less time.

Finally, among respondents who applied to more than one program and received at least one offer, 30% modified their ROL at some point between July 15 (the end of the application phase) and the time they completed the survey (half of them completed it before August 20). Among these students (panel C), 30% agreed that their perception of the universities was influenced by the early offers they received.

In sum, the survey results indicate that at the start of the procedure, many students have not yet formed preferences over the programs. Moreover, students tend to invest more time learning about universities from which they receive an offer than about others, and early offers influence their perceptions of the programs.

TABLE 5
EVIDENCE FROM A SURVEY ON STUDENTS IN THE DoSV FOR 2015–16

	Number of Valid Answers (1)	Agree/Yes (%) (2)
A. Applied to >1 Program (<i>N</i> = 4,994)		
At the time of application, I did not have a clear ranking because I still needed to collect information in order to rank my applications according to my preferences ^a	4,573	30
Getting to a ranking was very difficult, and I wanted to postpone this decision for as long as possible ^a	4,555	25
Received at least one offer	4,775	84
B. Applied to >1 Program and Received ≥1 Offer (<i>N</i> = 3,999)		
When comparing the universities that have made you an offer with universities that have not, can it then be said that: On average, I spend more time collecting information on the universities that have made me an offer ^a	3,251	61
On average, I spend the same amount of time collecting information on the universities that have made me an offer ^a	3,251	29
On average, I spend less time collecting information on the universities that have made me an offer ^a	3,251	10
Did your ranking change between the beginning of the procedure on July 15 and now?	3,552	30
C. Applied to >1 Program, Received ≥1 Offer, and Reranked Programs (<i>N</i> = 1,072)		
I have received some early offers that have changed my perception of the universities ^a	1,020	30

NOTE.—This table is based on the data from an online survey that was conducted between July 27 and October 10, 2015. The different panels correspond to different subgroups of respondents. Column 1 indicates the number of valid answers for each question, i.e., the number of participants who did not choose the option “I do not want to answer this question.” Among those who answered the question, col. 2 reports either the fraction of participants who responded yes (if the question requires a dichotomous yes/no answer) or the fraction who responded that they agree or strongly agree with the statement (if the question uses a five-point Likert scale).

^a These survey questions are originally based on a five-point Likert scale.

IV. A University Admissions Model with Preference Discovery

Here we develop a model of university admissions. Consistent with the findings from the survey in section III.B, it allows students’ preference

discovery through costly learning about programs. The sequential arrival of early offers under a multioffer mechanism such as the DoSV influences a student's learning behavior and leads to the documented early-offer effect. Further, we compare the DoSV with the canonical single-offer DA and a new mechanism—BM-DA—that improves on the DoSV. Recall that students apply to university programs and that each program admits students.

A. *The Mechanisms*

We start by formally defining the mechanisms. A multioffer DA (M-DA) mechanism, of which the DoSV is an example, is defined by the following stages:

Stage 1 (Applications).—Through a clearinghouse, students apply to a set of programs without committing to an ROL of these programs.

Stage 2 (Ranking applicants).—Every program ranks the students who have applied to it and submits the ranking to the clearinghouse.

Stage 3 (Multioffer stage with continuous offers and rejections).—Through the clearinghouse, each program extends admission offers to its top-ranked applicants up to its capacity. Students can reject any offers that they have received. Whenever a program's offer is rejected by a student, the clearinghouse automatically makes a new offer to the top-ranked applicant among those who have not yet received its offer.

Stage 4 (Final ROLs).—On a prespecified date, every student commits to an ROL of the programs that she has applied to (and has not rejected offers from).

Stage 5 (Final match).—With the rankings from students and programs as well as the remaining seats at each program, the clearinghouse runs the program-proposing GS algorithm and finalizes the matching.

In contrast to the M-DA, the program-proposing DA mechanism asks students to commit to an ROL of programs in stage 1 and proceeds to stages 2 and 5, while skipping all others. Note that we distinguish between “mechanism” and “algorithm,” although the literature often uses the two terms interchangeably. The GS algorithm is a set of computer codes that calculate a matching from information on students and programs, while a mechanism in our setting includes an algorithm and how relevant information is collected from students and programs (app. D.1).

A potential concern about the M-DA mechanism is that the sequential arrival of offers may lead to students accepting suboptimal offers. We thus propose the batched multioffer DA (BM-DA) mechanism that has a common date for every program to send out its early offers. Specifically, it replaces stage 3 of the M-DA with stage 3' while keeping everything else the same:

Stage 3' (Multioffer stage with offers batched on a common date).—On a prespecified date, every program extends admission offers through the

clearinghouse to its top-ranked applicants up to its capacity. Students can reject any offers that they have received.³³

To summarize, we distinguish between mechanisms under which a student can hold multiple offers for some time and those under which she receives at most one single offer. Accordingly, the DA is a single-offer mechanism, while the M-DA and the BM-DA allow multiple offers, with a more structured offer arrival under the BM-DA. Note that the multioffer mechanisms rely on the program-proposing GS algorithm instead of the student-proposing version. The latter cannot be used in a multioffer mechanism because the student-proposing GS algorithm allows each program only to tentatively accept a student before the algorithm concludes and thus does not generate offers until the very end.

B. Model Setup and Welfare Comparison

In the model, there is a student in a finite time horizon of $J + 2$ ($J \geq 2$) periods, $t \in \{0, 1, \dots, J, J + 1\}$.³⁴ We focus on two time points in each period, the beginning and the end, implying that actions can be taken by the student during a period. There is no time discounting.

At the beginning of period 0, the student has applied to a set of programs, \mathcal{J} ($|\mathcal{J}| = J$). In period 0, the student has the belief that she will receive an offer from $j \in \mathcal{J}$ in period $J + 1$ with probability $p_j^0 \in (0, 1)$. Let $p^0 \equiv (p_1^0, \dots, p_J^0)$. As we clarify shortly, the student does not update her belief about offer probabilities unless she receives an offer before period $J + 1$. Admission decisions are independent across programs.

Program j 's quality is a random variable, X_j , whose distribution function is G_j . Random variables X_j and $X_{j'}$ are independent for any $j \neq j'$. Let $G = \prod_j G_j$ represent the joint distribution.

The student's valuation of program j is $u(X_j)$, where $u(X_j)$ is weakly concave, implying that she can be risk neutral or risk averse. She has access to an outside option that can be an admission offer from another program outside of the system. This outside option gives her zero utility and is forfeited if she accepts an offer from any $j \in \mathcal{J}$.

At the beginning of period 0, the student knows only the distribution of X_j , G_j . She can pay a cost, c , to learn a program's (realized) quality and thus discover her preferences. As we specify below, the student's preference discovery is sequential. That is, she decides which program to learn about

³³ In some cases, it can be helpful to have multiple rounds of offers. We provide a discussion in sec. V.

³⁴ Unlike the matching literature, we study a single-agent model to highlight the learning dynamics in preference discovery. The comparison among the mechanisms therefore ignores equilibrium effects that other students' behavior may cause. Our simulations in sec. IV.D provide some idea about the magnitude of these equilibrium effects and reveal that the results from the single-agent problem still hold.

first, whether to learn about another program (conditional on what has been learned), and so on and so forth.

REMARK 1. Our model deviates from the common assumption in the matching literature that agents know their own ordinal preferences (Roth and Sotomayor 1990; Bogomolnaia and Moulin 2001). Under this assumption, appendix D.2 shows that the early-offer effect documented in section II cannot emerge. Intuitively, conditional on a student's applications, \mathcal{J} , knowing her own ordinal preferences is sufficient for her to make an optimal decision.

Periods 1 to J constitute the multioffer stage in the M-DA and BM-DA mechanisms. During this stage, there can be at most one early offer arriving at the beginning of each period. No offer can be rescinded. Let $O = (O_1, \dots, O_J) \in \{1, \dots, J, J+1\}^J$ represent an offer arrival, such that $O_j = t$ for $1 \leq t \leq J$ if the student receives an offer from program j in period t and that $O_j = J+1$ if she does not receive an offer from program j before period $J+1$. Note that $O_j = J+1$ does not imply that she will necessarily receive an offer from j in period $J+1$. There is no early offer in period 0.

Timeline under the DA.—The student is required to rank the J programs at the end of period 0 before any possible early offers. Therefore, her preference discovery activities, if any, take place in period 0.

Timeline under the BM-DA.—The student must rank the J programs at the end of period J . At the beginning of period J , she observes all early offers given O . The offer probabilities are updated from p^0 to $p^J(O)$, such that (i) $p_j^J(O) = p_j^0$ if $O_j = J+1$ and (ii) $p_j^J(O) = 1$ if $O_j < J+1$. She makes a decision on preference discovery given $p^J(O)$.

Timeline under the M-DA.—The student must submit an ROL of the J programs at the end of period J , but her learning decision is made myopically period by period as follows. Define $\underline{t} = \min\{\min\{O\}, J\}$, such that \underline{t} represents the period when her first early offer arrives, if any, or period J if there is no early offer. The student does not learn anything in period t for all $t < \underline{t}$. At the beginning of period t , for $\underline{t} \leq t \leq J$, the student updates her offer probabilities according to offer arrival O from $p^{(t-1)}(O)$ to $p^t(O) = (p_1^t(O), \dots, p_J^t(O))$, such that (i) $p_j^t(O) = p_j^0$ if $O_j > t$ (i.e., no offer from j up to period t) and (ii) $p_j^t(O) = 1$ if $O_j \leq t$ (i.e., an offer from j arrives in period t or earlier). Conditional on $p^t(O)$ and what the student has learned up to period $(t-1)$, she decides whether to learn more, assuming (incorrectly) that there will be no more offers before period $(J+1)$. If $\min\{O\} = J+1$ (i.e., there is no early offer at all), she makes a decision on preference discovery in period J given that $p^J(O) = p^0$.

We assume that under the M-DA the student ignores the possibility of receiving more early offers in later periods. In practice, some students under the M-DA behave in this manner, as we have seen in our data. If

all students are fully rational, they would not make a learning decision until all early offers have been received, making the DoSV (an application of the M-DA) equivalent to the BM-DA (conditional on offer arrival). In this case, there would be no difference between the first and later early offers, provided that offers arrive randomly. However, section II.C shows that the first early offer has a larger effect than later early offers.

REMARK 2. As shown below, depending on the mechanism in place, the student endogenously discovers her preferences over the programs. Given her information set at the time of ROL finalization, each mechanism has the same properties as the program-proposing DA: it is stable with respect to reported ordinal preferences and not strategy-proof for students or programs.³⁵

1. Learning Strategy in a Given Period

We now consider the student's learning strategy in a given period with up-to-date offer probabilities $p \in (0, 1]^J$. This is precisely what she does under the DA or BM-DA, because preference discovery happens only in one period. We will extend the derivation to discovery in multiple periods under the M-DA.

Let $\mathcal{U} \subseteq \mathcal{J}$ represent the set of programs whose qualities have not yet been learned by the student, and let $\bar{\mathcal{U}} = \mathcal{J} \setminus \mathcal{U}$. Also let $l(\bar{\mathcal{U}}) = \{(j, x_j)_{j \in \bar{\mathcal{U}}}\} \in (\mathcal{J} \times \mathbb{R})^{|\bar{\mathcal{U}}|}$ represent the realized program qualities already learned by the student, paired with their identities.

Given a state (\mathcal{U}, l) and offer probabilities p , the student's learning strategy is given by $\psi(\mathcal{U}, l|p) \in \{0\} \cup \mathcal{U}$, such that $\psi = 0$ when the student stops learning and $\psi \in \mathcal{U}$ indicates which program to learn next. With this definition, we consider only pure strategies.³⁶

Given (\mathcal{U}, l, p) , if $\psi = 0$, the student then submits an ROL. The expected utility from an ROL is calculated as follows. For $j \in \bar{\mathcal{U}}$, define $v_j = \max\{0, u(x_j)\}$, in which x_j is the realized value of X_j ; for $j \in \mathcal{U}$, $v_j = \max\{0, \int u(x) dG_j(x)\}$. Recall that the outside option yields zero utility. Given her information, it is optimal for the student to submit a truthful ROL of all acceptable programs (i.e., $v_j > 0$).³⁷ Not taking into account the learning costs that are already sunk, the expected utility

³⁵ Under the assumption that student preferences evolve over time exogenously, Narita (2018) investigates the properties of alternative mechanisms.

³⁶ A mixed strategy implies that in some state, the student is indifferent between two different actions. This may happen for a measure zero set of (\mathcal{U}, l) when G_j admits a continuous density function for all j .

³⁷ We assume that including a program in a submitted ROL is a commitment to accepting the program's offer if it is the highest-ranked offer. Equivalently, we may assume that there is a cost to reject an unacceptable offer that is included in the student's ROL.

from this ROL is given by $v(\mathcal{U}, l, p) = p_{j_l} v_{j_l} + \sum_{k=2}^J p_{j_k} v_{j_k} \prod_{k'=1}^{k-1} (1 - p_{j_{k'}})$, where $v_{j_k} \geq v_{j_{k'}}$ whenever $k < k'$.

If the student adopts strategy ψ , her expected utility given (\mathcal{U}, l, p) is given by

$$V(\mathcal{U}, l, p|\psi) = \mathbf{1}_{(\psi(\mathcal{U}, l|p)=0)} \times v(\mathcal{U}, l, p) + \sum_{j \in \mathcal{U}} \mathbf{1}_{(\psi(\mathcal{U}, l|p)=j)} \left(\int V(\mathcal{U} \setminus \{j\}, l \cup \{(j, x_j)\}, p|\psi) dG_j(x_j) - c \right),$$

where $\mathbf{1}_{(\cdot)}$ is an indicator function. Thus, $V(\mathcal{J}, \emptyset, p|\psi)$ represents the student's expected payoff from adopting strategy ψ before she starts discovering her preferences.

An optimal strategy, ψ^* , solves the following problem for every state (\mathcal{U}, l) given p :

$$V(\mathcal{U}, l, p|\psi^*) = \max \left\{ v(\mathcal{U}, l, p), \max_{j \in \mathcal{U}} \left\{ \int V(\mathcal{U} \setminus \{j\}, l \cup \{(j, x_j)\}, p|\psi^*) dG_j(x_j) - c \right\} \right\}.$$

In other words, $\psi^*(\mathcal{U}, l|p) = 0$ if $v(\mathcal{U}, l, p) \geq \int V(\mathcal{U} \setminus \{j\}, l \cup \{(j, x_j)\}, p|\psi^*) dG_j(x_j) - c$ for all $j \in \mathcal{U}$; otherwise, $\psi^*(\mathcal{U}, l|p) = \arg \max_{j \in \mathcal{U}} \left\{ \int V(\mathcal{U} \setminus \{j\}, l \cup \{(j, x_j)\}, p|\psi^*) dG_j(x_j) \right\}$. By definition, $V(\mathcal{J}, \emptyset, p|\psi^*) \geq V(\mathcal{J}, \emptyset, p|\psi)$ for all ψ .

Preference discovery under the DA.—The student discovers her preferences over the programs in period 0. That is, she starts learning given $(\mathcal{J}, \emptyset, p^0)$. Let $\psi^{DA}(\cdot, \cdot | p^0)$ represent an optimal strategy.

Preference discovery under the BM-DA.—The student learns about her preferences in period J after receiving all early offers as prescribed by offer arrival O . Hence, she starts learning given $(\mathcal{J}, \emptyset, p^j(O))$. Let $\psi^B(\cdot, \cdot | p^j(O))$ represent an optimal strategy.

2. Preference Discovery under the M-DA

We now extend the analysis to the M-DA under which the student's preference discovery can be updated upon the arrival of an early offer. Recall that $\underline{t} = \min\{\min\{O\}, J\}$. In each period $t = \underline{t}, \dots, J$, given $p^t(O)$, the student has a myopic strategy, $\psi^t(\cdot, \cdot | p^t(O))$, leading to (subjective) expected utility $V^t(\mathcal{U}^{(t-1)}, l^{(t-1)}, p^t(O)|\psi^t(\cdot, \cdot | p^t(O)))$, where $\mathcal{U}^{(t-1)}$ and $l^{(t-1)}$ represent the learning outcomes at the end of period $(t-1)$ and $\mathcal{U}^{(\underline{t}-1)} = \mathcal{J}$ and $l^{(\underline{t}-1)} = \emptyset$. Hence, $\psi^t \in \arg \max_{\psi} V^t(\mathcal{U}^{(t-1)}, l^{(t-1)}, p^t(O)|\psi)$ for $t = \underline{t}, \dots, J$.

The M-DA thus leads to a sequence of learning strategies, $\{\psi^t\}_{t=\underline{t}}^J$. Importantly, the learning technology is such that the student cannot unlearn what has been learned. We show in appendix D.4 that under an assumption of no learning on off-equilibrium paths (assumption D1), for

any state, there is at most one strategy in the sequence $\{\psi^i\}_{i=\underline{t}}^J$ requiring the student to learn more (lemma D1). Therefore, we can define $\psi^M(\mathcal{U}, l|p^J(O)) = \max_{i \in \{\underline{t}, \dots, J\}} \psi^i(\mathcal{U}, l|p^i(O))$ for all (\mathcal{U}, l) . Under lemma D1, $\psi^M(\mathcal{U}, l|p^J(O))$ is equivalent to applying $\{\psi^i\}_{i=\underline{t}}^J$ sequentially.

3. Welfare Dominance of the BM-DA over the DA and M-DA

For each mechanism, we study the student's welfare at the beginning of period 0 (i.e., before any preference discovery) conditional on offer arrival O . Thus, we do not need to specify a distribution for O .

Under the DA, the expected utility at the beginning of period 0 is evaluated at $p^J(O)$, resulting in $V(\mathcal{J}, \emptyset, p^J(O)|\psi^{DA}(\cdot, \cdot|p^0))$. Notice that the learning strategy is determined based on p^0 without knowing the offer arrival O , which is a source of inefficiency. Under the M-DA, it is $V(\mathcal{J}, \emptyset, p^J(O)|\psi^M(\cdot, \cdot|p^J(O)))$, and under the BM-DA it is $V(\mathcal{J}, \emptyset, p^J(O)|\psi^B(\cdot, \cdot|p^J(O)))$. We then have the following result.

PROPOSITION 1. Conditional on the offer arrival O , in terms of the student's expected utility at the beginning of period 0,

- (i) the BM-DA mechanism dominates both the M-DA and the DA mechanisms:

$$\begin{aligned} V(\mathcal{J}, \emptyset, p^J(O)|\psi^B(\cdot, \cdot|p^J(O))) &\geq V(\mathcal{J}, \emptyset, p^J(O)|\psi^M(\cdot, \cdot|p^J(O))), \\ V(\mathcal{J}, \emptyset, p^J(O)|\psi^B(\cdot, \cdot|p^J(O))) &\geq V(\mathcal{J}, \emptyset, p^J(O)|\psi^{DA}(\cdot, \cdot|p^0)), \end{aligned}$$

with each inequality being strict for some $(\mathcal{J}, p^0, O, F, c)$, and

- (ii) the welfare ranking between the M-DA and the DA is ambiguous.

For an arbitrary distribution of O , part i still holds in terms of the student's unconditional expected utility in period 0 because it is satisfied for all possible O .

C. Costly Preference Discovery and Early-Offer Effect: A Numerical Example

Our empirical analysis documents that an early offer increases the probability that the offer is accepted and that the effect is larger for the first offer. We now present a numerical example to highlight how these effects can be caused by the student's learning about program qualities.³⁸

³⁸ A working paper version of this article, Grenet, He, and Kübler (2021), considers a more general setting. It implies that the numerical example here is not a knife-edge case. We also note that the positive early-offer effect on offer acceptance does not always emerge. Appendix D.3 provides an example in which an early offer from a program reduces the

Moreover, we clarify a disadvantage of the M-DA: if a low-quality program makes an offer first, this may suboptimally increase the probability that the student will accept its offer. We show that this inefficiency is avoided by the BM-DA.

The student is risk neutral, $u(X) = X$. There are two programs, $\mathcal{J} = \{1, 2\}$. The quality of program j is an independent draw from $\text{Uniform}(\mu_j - 0.5, \mu_j + 0.5)$, with $\mu_1 = 1/16$ and $\mu_2 = 1/32$. As seen in period $t = 0$, the probability of receiving an offer in period $t = J + 1$ from program 1 is $p_1^0 = 9/16$ and that from program 2 is $p_2^0 = 9/16$. The cost of learning a program's quality is $c = 1/20$. Everything else is the same as in section IV.B. For instance, there is an outside option whose value is zero.

There are five different types of offer arrivals, O^\emptyset (no early offer), $O^{(1)}$ (an early offer from program 1), $O^{(2)}$ (an early offer from program 2), $O^{(1,2)}$ (early offers from program 1 and then program 2), and $O^{(2,1)}$ (early offers from program 2 and then program 1). Below, we investigate optimal learning (sec. IV.C.1), ranking behavior (sec. IV.C.2), and the efficiency of the mechanisms (sec. IV.C.3).

1. Learning Behavior

The student's learning strategy under a given mechanism has two parts: which program to learn first and, depending on what has been learned, whether to learn about the other program. We solve her optimization problem by backward induction.

Suppose that the offer arrival is O^\emptyset (no early offer). If the student first learns X_1 , the decision of whether to learn X_2 is depicted in figure 5A. When the realization of X_1 , x_1 , is below the threshold (the dashed vertical line on the right), $X_1^*(p_1^0, p_2^0) = \mu_2 + 1/2 - ((2c - (1 - p_1^0)p_2^0(\mu_2 - 1/2))/p_1^0 p_2^0)^{1/2} \approx 0.150$, the student learns X_2 ; otherwise, she stops learning. Intuitively, with a higher value of x_1 , the incentive to learn X_2 is lower because she can stop learning and top rank program 1 to guarantee herself an expected payoff of $p_1^0 x_1 + (1 - p_1^0)p_2^0 \mu_2$. Similar thresholds can be derived when the student learns X_2 first and/or when there are one or two early offers (for detailed derivations, see Grenet, He, and Kübler 2021).

We now consider which program to learn first. For $j, j' \in \{1, 2\}$, let X_j -then- $X_{j'}$ denote that the student first learns X_j and then, depending on the realization of X_j , possibly learns $X_{j'}$. It involves pairwise welfare comparisons among X_1 -then- X_2 , X_2 -then- X_1 , and no learning. Given O^\emptyset and the parameter values, we can show that it is optimal to learn X_1 first, as depicted in figure 5A.

likelihood that the offer is accepted. Besides, there is no first-early-offer effect. The main feature of that example is that the student never has any incentive to learn about one of the two programs.

RESULT 1. Suppose that the offer arrival is O . The student's optimal learning sequence is as follows:

- (i) if $O = O^\emptyset$, X_1 -then- X_2 under the M-DA or the BM-DA;
- (ii) if $O = O^{\{1\}}$, X_1 -then- X_2 under the M-DA or the BM-DA;
- (iii) if $O = O^{\{2\}}$, X_2 -then- X_1 under the M-DA or the BM-DA;
- (iv) if $O = O^{\{1,2\}}$, X_1 -then- X_2 under the M-DA or the BM-DA;
- (v) if $O = O^{\{2,1\}}$, X_2 -then- X_1 under the M-DA but X_1 -then- X_2 under the BM-DA.
- (vi) Under the DA, since the student has to carry out preference discovery before receiving any early offer, the optimal sequence is always X_1 -then- X_2 .




FIG. 5.—Learning and ranking behaviors with or without an early offer from program 1. This figure applies to both the M-DA and the BM-DA. We set $\mu_1 = 1/16$, $\mu_2 = 1/32$, $p_1^0 = 9/16$, $p_2^0 = 9/16$, and $c = 1/20$. In A, the student receives no early offer. She learns X_1 first and continues to learn X_2 if and only if the realization of X_1 , x_1 , is below $X_1^*(p_1^0, p_2^0) = \mu_2 + 1/2 - ((2c - (1 - p_1^0)p_2^0(\mu_2 - 1/2)^2)/p_1^0 p_2^0)^{1/2}$. If $x_1 \in [X_1^*(p_1^0, p_2^0), \mu_1 + 1/2]$, the student submits ROL 1-2. If $x_1 \in [0, X_1^*(p_1^0, p_2^0)]$, she will learn X_2 and submits a truthful ROL among 1-2, 2-1, and 1-0. If $x_1 \in [\mu_1 - (1/2), 0]$, she will learn X_2 and submits a truthful ROL, either 2-0 or 0-0 (an empty ROL). In B, the student receives an early offer from program 1, and it is still optimal to learn X_1 first. The probability of learning X_2 decreases because the threshold drops from $X_1^*(p_1^0, p_2^0)$ to $X_1^*(1, p_2^0) = \mu_2 + 1/2 - (2c/p_2^0)^{1/2}$. The probability

TABLE 6
LEARNING, RANKING, AND WELFARE GIVEN OFFER ARRIVAL

Mechanism	O^\emptyset : No Early Offer (1)	$O^{(1)}$: Early Offer from Program 1 (2)	$O^{(2)}$: Early Offer from Program 2 (3)	$O^{(1,2)}$: Early Offers from Program 1 and Then 2 (4)	$O^{(2,1)}$: Early Offers from Program 2 and Then 1 (5)
A. Discovery Probabilities ($\times 100$)					
DA:					
Pr(learn X_1)	100	100	100	100	100
Pr(learn X_2)	58.8	58.8	58.8	58.8	58.8
M-DA:					
Pr(learn X_1)	100	100	61.0	100	71.5
Pr(learn X_2)	58.8	54.7	100	65.3	100.0
BM-DA:					
Pr(learn X_1)	100	100	88.6	100	100
Pr(learn X_2)	58.8	54.7	100	65.3	65.3
B. Probabilities of Top Ranking a Program ($\times 100$)					
DA:					
Pr(top rank program 1)	49.4	49.4	49.4	49.4	49.4
Pr(top rank program 2)	30.1	30.1	30.1	30.1	30.1
M-DA:					
Pr(top rank program 1)	49.4	51.0	33.3	47.1	37.2
Pr(top rank program 2)	30.1	28.5	46.2	32.4	42.3
BM-DA:					
Pr(top rank program 1)	49.4	51.0	33.3	47.1	47.1
Pr(top rank program 2)	30.1	28.5	46.2	32.4	32.4
C. Welfare Conditional on Offer Arrival ($\times 1,000$)					
DA	57.7	122.4	95.1	156.3	156.3
M-DA	57.7	122.6	112.0	157.0	153.9
BM-DA	57.7	122.6	112.0	157.0	157.0

NOTE.—These results are calculated for $\mu_1 = 1/16$, $\mu_2 = 1/32$, $p_1^0 = 9/16$, $p_2^0 = 9/16$, and $c = 1/20$.

For the M-DA, result 1 considers the learning sequence in the period when the first offer arrives, provided that $O \neq O^\emptyset$. It shows that offer arrival determines the optimal learning sequence. In turn, it affects the probability of a program’s quality being learned; when the optimal sequence is X_j —then— X_j , X_j is learned for certain while X_i is learned only with some probability. This is illustrated in panel A of table 6. In column 1, given O^\emptyset , the probability of X_1 being learned is one under the M-DA or BM-DA, but it decreases to 61.0% under the M-DA or BM-DA when there is

of top ranking program 1 increases: for $x_1 \in [X_1^*(1, p_2^0), X_1^*(p_1^0, p_2^0)]$, she now always submits ROL 1–2 instead of one of the ROLs 1–2, 2–1, and 1–0.

an early offer from program 2 (col. 3). This is because the early offer changes the optimal learning sequence.

Even when the optimal sequence remains the same, the learning probability can change for the program that is learned second. In column 2 of table 6, when the student receives an early offer from program 1 and thus the offer arrival changes from O^\emptyset to $O^{(1)}$, the probability of X_2 being learned decreases from 58.8% to 54.7% under either the M-DA or the BM-DA. This is depicted in figure 5: the threshold of X_1 , X_1^* , below which the student does not learn X_2 is lower given $O^{(1)}$ (panel B) than given O^\emptyset (panel A).

2. Ranking Behavior under the M-DA

We now explore how early offers affect the student's ROL of the programs. Investigating the student's ranking behavior amounts to studying her offer acceptance because conditional on eventually receiving offers from both programs, the top-ranked program will be accepted by the student. This is similar to our empirical analysis of offer acceptance in which we focus on a student's ex post feasible programs that would make her an offer.

Figure 5A shows that given O^\emptyset and if $x_1 > X_1^*(p_1^0, p_2^0)$, the student stops learning and submits ROL 1–2. If $0 < x_1 \leq X_1^*(p_1^0, p_2^0)$, she learns X_2 and submits a truthful ROL, 1–2, 2–1, or 1–0. If $x_1 < 0$, she learns X_2 and submits a truthful ROL, 2–0 or 0–0 (an empty ROL). However, when the offer arrival becomes $O^{(1)}$ (panel B), for $X_1^*(1, p_2^0) < x_1 < X_1^*(p_1^0, p_2^0)$, the new optimal strategy is to stop learning and submit ROL 1–2, while there is no change for other values of x_1 . This increases the probability that program 1 is top ranked, from 49.4% to 51.0% under the M-DA as shown in columns 1 and 2 in panel B of table 6. The table provides more examples of this early-offer effect that are summarized in result 2.

RESULT 2. Receiving an offer from one of the programs, either as the first or as the second early offer, strictly increases the probability that this program is top ranked in the student's submitted ROL under the M-DA, *ceteris paribus*.

Result 2 describes the effect of receiving an early offer from j , holding everything else constant. For example, it calculates the difference in the probability of program 1 being top ranked between $O^{(1)}$ and O^\emptyset , between $O^{(1,2)}$ and $O^{(2)}$, or between $O^{(2,1)}$ and $O^{(2)}$.

When there are two early offers, arrival sequence matters under the M-DA. The following result shows the first-early-offer effect on ranking behavior under the M-DA.

RESULT 3. Suppose that there are two early offers. Under the M-DA, the probability of a program being top ranked in the student's submitted ROL is higher when its offer is the first one than when it is the second.

Result 3 focuses on the difference in the probability of program 1 being top ranked between $O^{(1,2)}$ and $O^{(2,1)}$; a similar comparison is also made

for program 2. For example, when we move from column 4 to column 5 in panel B of table 6, the early offer from program 2 becomes the first offer and the probability of program 2 being top ranked increases from 32.4% to 42.3% under the M-DA. Results 2 and 3 are consistent with our empirical findings in section II.

3. Welfare Dominance of the BM-DA and Welfare Loss under the M-DA

We now compare student welfare under the three mechanisms. Given the parameter values, we show that some of the weak inequalities in proposition 1 are strict. In particular, we highlight that the M-DA can result in a welfare loss when a low-quality offer arrives first.

RESULT 4. Suppose that the offer arrival is O . We have the following results on student welfare:

1. *Dominance of the BM-DA.* Conditional on O , the BM-DA dominates the DA, and the dominance is strict if $O \neq O^\varnothing$ (i.e., there is at least one early offer); the BM-DA also dominates the M-DA, and the dominance is strict when $O = O^{(2,1)}$ (i.e., early offers from program 2 and then program 1).
2. *Comparison between the DA and the M-DA.* Conditional on O , the M-DA is equivalent to the DA if $O = O^\varnothing$; the DA strictly dominates the M-DA if $O = O^{(2,1)}$; and the M-DA strictly dominates the DA if $O \neq O^\varnothing$ and $O \neq O^{(2,1)}$.

Result 4 is illustrated in panel C of table 6. Specifically, given $O = O^{(2,1)}$ (col. 5), the M-DA leads to a welfare loss relative to both the DA and the BM-DA. Under the M-DA, the student's optimal learning sequence is X_2 -then- X_1 , starting with the ex ante low-quality program (panel A). As result 3 shows, the student has a higher probability of top ranking program 2 (panel B), often without learning X_1 . Hence, the student misses some chances of being assigned to program 1. This inefficiency is avoided by the BM-DA under which the optimal learning sequence is X_1 -then- X_2 .

D. Simulation Results

To gauge the relative performance of the DA, M-DA, and BM-DA mechanisms in a real-life setting, we calibrate a set of simulations with the same data used in section II. Owing to the data limitations and the lack of a fully fledged empirical model, certain simplifying assumptions are needed in the simulations. We highlight these assumptions below, while the details of the simulation procedure are in appendix E.

1. Setup

We construct a stylized market based on the DoSV data. As in section II, we use the 21,711 students who applied to at least two feasible programs and accepted an offer. In every simulation, student i always applies to the same set of programs, \mathcal{A}_i , including those in her initial ROL and an outside option.³⁹ A program's capacity is set equal to the number of students in the simulation sample who accepted its offer in reality. Each program ranks its applicants by the average of their *Abitur* percentile rank (between zero and one) and a program-specific Uniform(0, 1) random variable.

The timeline is as follows: (i) under the DA, students submit their ROL before receiving any offer; (ii) under the M-DA, each program sends out a single batch of early offers at a random time to its highest-ranked applicants up to its capacity and the ROL submission is after the arrival of early offers; (iii) under the BM-DA, the timing is the same as under the M-DA except that all early offers are sent out on the same date. Under any mechanism, the matching is determined by the program-proposing GS algorithm, using as inputs the submitted ROLs and the programs' capacities and rankings of students.

2. Preference Discovery

As in our theoretical model, students are uncertain about their preferences over programs and can learn their quality only at a cost. If student i pays the cost, she discovers her true utility from program j , U_{ij}^{FullInfo} ; otherwise, she values j at U_{ij}^{NoInfo} . At the time of submitting her ROL, student i 's perceived utility from program j under mechanism m is specified as

$$U_{ij}^m = \lambda_{ij}^m \cdot U_{ij}^{\text{FullInfo}} + (1 - \lambda_{ij}^m) U_{ij}^{\text{NoInfo}} \quad \forall i, j \in \mathcal{A}_i, \quad (3)$$

where $\lambda_{ij}^m \in \{0, 1\}$ is an indicator for whether i has learned her true utility from j .

Further, $U_{ij}^{\text{FullInfo}} = V_{ij}^{\text{FullInfo}} + \epsilon_{ij}^{\text{FullInfo}}$, where V_{ij}^{FullInfo} depends on observable student-program-specific characteristics and $\epsilon_{ij}^{\text{FullInfo}}$ is i.i.d. type I extreme value distributed. To calculate V_{ij}^{FullInfo} , we rely on the same sample as in section II.⁴⁰ Under the assumption that a student always learns her preference for her first early offer, V_{ij}^{FullInfo} is computed by imposing that j

³⁹ This outside option accounts for the possibility that students may have applied to programs outside the DoSV procedure and devoted some time to discovering their quality. Since all students in our sample accepted an offer from the clearinghouse, we make the simplifying assumption that this outside option is never feasible.

⁴⁰ We estimate a rank-ordered logit model that is an augmented version of the specification in col. 3 of table 3, in which we fully interact the early-offer and first-early-offer dummies with university fixed effects, field-of-study fixed effects, distance, distance squared, and a dummy for whether the program is in the student's region. This model is estimated using the information extracted from a student's final ROL.

is her first early offer. Similarly, we assume that $U_{ij}^{\text{NoInfo}} = V_{ij}^{\text{NoInfo}} + \epsilon_{ij}^{\text{NoInfo}}$, where V_{ij}^{NoInfo} is drawn from a normal distribution with mean V_{ij}^{FullInfo} and $\epsilon_{ij}^{\text{NoInfo}}$ is a type I extreme value such that $\epsilon_{ij}^{\text{NoInfo}} \perp \epsilon_{ij}^{\text{FullInfo}}$.

Lacking an estimate of learning costs, we assume that under any mechanism, student i discovers her true preferences for half of the programs in \mathcal{A}_i . Further deviating from the endogenous discovery in section IV.B, we impose the following assumptions. Under the DA, the programs whose qualities are learned by i (i.e., $\lambda_{ij}^{\text{DA}} = 1$) are randomly selected. Under the M-DA, we emulate the myopic learning behavior in section IV.B by assuming that i always learns the quality of her first early offer and then alternates between (i) learning the quality of a randomly chosen program from the ones in \mathcal{A}_i that she has not learned yet and (ii) learning the quality of her early offers (if any) in their arrival order. Under the BM-DA, early offers are learned before other programs.

3. Results

The simulations generate students' submitted ROLs and matching outcomes under the three mechanisms in 10,000 samples. As a benchmark, we also simulate the matching that would have been observed under full information.

Figure 6A contrasts the mechanisms by students' preference discovery. The simulations confirm that in terms of discovery, (i) the M-DA and BM-DA mechanisms are both more efficient than the DA and (ii) the BM-DA slightly outperforms the M-DA. On average, across the simulation samples, 44.0% of students rank ex post feasible programs in the order of their full information preferences under the BM-DA, relative to 43.3% and 40.6% under the M-DA and the DA, respectively.⁴¹

As for student welfare, we compute each student's expected utility as the average utility (U_{ij}^{FullInfo}) from her matches across the simulation samples. We then perform pairwise comparisons of the mechanisms based on the shares of students whose expected utility is (i) strictly higher under one mechanism than under the other, (ii) strictly lower, or (iii) equal. These comparisons incorporate equilibrium effects since a student's match depends not only on her submitted ROL but also on other students' behavior. As shown in figure 6B, the majority of students (over 70%) are strictly better off under the M-DA and the BM-DA than under the DA. Moreover, the BM-DA outperforms the M-DA, as 41.7% of students are strictly better off under the BM-DA than under the M-DA, while only 37.1% are worse off.

⁴¹ The comparisons are restricted to ex post feasible programs since it is payoff-irrelevant for a student to rank an infeasible program arbitrarily.




FIG. 6.—Comparing the DA, M-DA, and BM-DA mechanisms: simulation results. Using the simulations detailed in appendix E, this figure compares the DA, M-DA, and BM-DA mechanisms along two dimensions: students' submitted ROLs (A) and their expected utility of the matching outcome (B). There are 21,711 students. A, Comparison of the share of students who, under each mechanism, rank the ex post feasible programs to which they applied in the order of their full-information preferences, averaged across 10,000 simulation samples. B, Pairwise comparisons of mechanisms based on the shares of students whose expected utility is (i) strictly higher under the first mechanism than under the second, (ii) strictly lower, or (iii) equal. A student's expected utility is the average of the full-information utility of her matches across the 10,000 simulation samples.

Overall, these simulations provide suggestive evidence that the M-DA and the BM-DA can generate sizable welfare gains relative to the DA. Furthermore, the batching of early offers in the BM-DA mechanism is found to yield a small yet meaningful improvement on the M-DA.

V. Concluding Remarks

In recent years, matching markets have become more centralized and increasingly rely on single-offer mechanisms. This trend has been especially prominent in school choice and university admissions where students are required to rank a large number of schools or universities before receiving an offer. Theoretical justifications are usually based on the assumption that agents know their own preferences upon participating in a matching market.

Relying on a unique data set from Germany's university admissions, we identify a quasi experiment in which the arrival time of admission offers is as good as random. We show that a student is more likely to accept an early offer relative to later offers, despite the fact that offers do not expire until the end of the admission procedure. This finding cannot be reconciled with the assumption that students have full information on their own preferences over university programs. Instead, it is consistent with students' discovering their preferences at a cost, which is corroborated by survey evidence.

These results highlight the advantages of dynamic multioffer mechanisms. Before students are required to rank university programs, they are allowed to hold more than one offer in the multioffer phase, which helps them optimize their preference discovery. Our proposed BM-DA mechanism, which is an improvement over the multioffer mechanism used in Germany's university admissions, captures such efficiency gains and dominates the DA, the most commonly used single-offer mechanism in real-world matching markets.

The BM-DA mechanism can be implemented through an online clearinghouse, similar to the DoSV in Germany and *Parcoursup* in France for university admissions. When implementing the BM-DA mechanism in practice, it is important to pay attention to the details of its design. For example, the choice of the number of rounds with offers in the multioffer phase should be guided by the specific context. The more time-consuming the preference discovery is, the more time students should have between rounds, *ceteris paribus*, which implies fewer rounds of offers if the total length of the multioffer phase is fixed. Moreover, highly homogeneous preferences of both students and programs imply that offers in each round are concentrated on a small group of students. Having more rounds of offers may benefit students who are ranked low by the programs. We leave the formal analysis of such considerations to future research.

References

- Abdulkadiroğlu, Atila, Joshua D. Angrist, Yusuke Narita, and Parag A. Pathak. 2017. "Research Design Meets Market Design: Using Centralized Assignment for Impact Evaluation." *Econometrica* 85 (5): 1373–432.
- Abdulkadiroğlu, Atila, Yeon-Koo Che, Parag A. Pathak, Alvin E. Roth, and Olivier Tercieux. 2020. "Efficiency, Justified Envy, and Incentives in Priority-Based Matching." *A.E.R.: Insights* 2 (4): 425–42.
- Abdulkadiroğlu, Atila, Parag A. Pathak, and Alvin E. Roth. 2005. "The New York City High School Match." *A.E.R.* 95 (2): 364–67.
- Abdulkadiroğlu, Atila, Parag A. Pathak, Alvin E. Roth, and Tayfun Sönmez. 2005. "The Boston Public School Match." *A.E.R.* 95 (2): 368–71.
- Abdulkadiroğlu, Atila, and Tayfun Sönmez. 2003. "School Choice: A Mechanism Design Approach." *A.E.R.* 93 (3): 729–47.
- Alladi, Vinayak. 2018. "Sour Grapes in the Lab and Field: A Test of Access-Based Beliefs." Working paper, Univ. California, San Diego.
- Antler, Yair. 2019. "No One Likes to Be Second Choice." *Econ. J.* 129 (619): 1119–38.
- Artemov, Georgy. 2021. "Assignment Mechanisms: Common Preferences and Information Acquisition." *J. Econ. Theory* 198:105370.
- Artemov, Georgy, Yeon-Koo Che, and YingHua He. 2017. "Strategic 'Mistakes': Implications for Market Design Research." 2017. Working paper.
- Avery, Christopher, Andrew Fairbanks, and Richard J. Zeckhauser. 2003. *The Early Admissions Game: Joining the Elite*. Cambridge, MA: Harvard Univ. Press.
- Avery, Christopher N., Mark E. Glichman, Caroline M. Hoxby, and Andrew Metrick. 2013. "A Revealed Preference Ranking of U.S. Colleges and Universities." *Q.J.E.* 128 (1): 425–67.
- Avery, Christopher, and Jonathan Levin. 2010. "Early Admissions at Selective Colleges." *A.E.R.* 100 (5): 2125–56.
- Bergemann, Dirk, and Juuso Välimäki. 2002. "Information Acquisition and Efficient Mechanism Design." *Econometrica* 70 (3): 1007–33.
- Bó, Inació, and Rustamdjan Hakimov. 2018. "The Iterative Deferred Acceptance Mechanism." Working paper.
- . 2020. "Iterative versus Standard Deferred Acceptance: Experimental Evidence." *Econ. J.* 130 (626): 356–92.
- Bogomolnaia, Anna, and Hervé Moulin. 2001. "A New Solution to the Random Assignment Problem." *J. Econ. Theory* 100 (2): 295–328.
- Braun, Sebastian, Nadja Dwenger, and Dorothea Kübler. 2010. "Telling the Truth May Not Pay Off: An Empirical Study of Centralized University Admissions in Germany." *B.E. J. Econ. Analysis and Policy* 10 (1): 22.
- Braun, Sebastian, Nadja Dwenger, Dorothea Kübler, and Alexander Westkamp. 2014. "Implementing Quotas in University Admissions: An Experimental Analysis." *Games and Econ. Behavior* 85 (C): 232–51.
- Chen, Yan, and YingHua He. 2021a. "Information Acquisition and Provision in School Choice: A Theoretical Investigation." *Econ. Theory*.
- . 2021b. "Information Acquisition and Provision in School Choice: An Experimental Study." *J. Econ. Theory* 197:105345.
- Chen, Yan, and Tayfun Sönmez. 2006. "School Choice: An Experimental Study." *J. Econ. Theory* 127 (1): 202–31.
- Compte, Olivier, and Philippe Jehiel. 2007. "Auctions and Information Acquisition: Sealed Bid or Dynamic Formats?" *RAND J. Econ.* 38 (2): 355–72.
- Dizdar, Deniz, and Benny Moldovanu. 2016. "On the Importance of Uniform Sharing Rules for Efficient Matching." *J. Econ. Theory* 165 (C): 106–23.

- Doval, Laura. 2018. "Whether or Not to Open Pandora's Box." *J. Econ. Theory* 175 (C): 127–58.
- Dwenger, Nadja, Dorothea Kübler, and Georg Weizsäcker. 2018. "Flipping a Coin: Evidence from University Applications." *J. Public Econ.* 167:240–50.
- Dzyabura, Daria, and John R. Hauser. 2019. "Recommending Products When Consumers Learn Their Preferences." *Marketing Sci.* 38 (3): 417–41.
- Echenique, Federico, Alistair J. Wilson, and Leeat Yariv. 2016. "Clearinghouses for Two-Sided Matching: An Experimental Study." *Quantitative Econ.* 7 (2): 449–82.
- Echenique, Federico, and Leeat Yariv. 2013. "An Experimental Study of Decentralized Matching." Working paper, Dept. Econ., Princeton Univ., Princeton, NJ.
- Elster, Jon. 1983. *Sour Grapes: Studies in the Subversion of Rationality*. Cambridge: Cambridge Univ. Press.
- Evans, Jonathan St. B. T. 2008. "Dual-Processing Accounts of Reasoning, Judgment, and Social Cognition." *Ann. Rev. Psychology* 59:255–78.
- Fack, Gabrielle, Julien Grenet, and YingHua He. 2019. "Beyond Truth-Telling: Preference Estimation with Centralized School Choice and College Admissions." *A.E.R.* 109 (4): 1486–529.
- Fudenberg, Drew, and David K. Levine. 2006. "A Dual-Self Model of Impulse Control." *A.E.R.* 96 (5): 1449–76.
- Gale, David E., and Lloyd S. Shapley. 1962. "College Admissions and the Stability of Marriage." *American Math. Monthly* 69 (1): 9–15.
- Gong, Binglin, and Yingzhi Liang. 2020. "A Dynamic Matching Mechanism for College Admissions: Theory and Experiment." Working paper.
- Gossner, Olivier, Jakub Steiner, and Colin Stewart. 2021. "Attention Please!" *Econometrica* 89 (4): 1717–51.
- Grenet, Julien, YingHua He, and Dorothea Kübler. 2021. "Decentralizing Centralized Matching Markets: Implications from Early Offers in University Admissions." Working paper.
- Haeringer, Guillaume, and Vincent Iehlé. 2021. "Gradual College Admission." *J. Econ. Theory* 198:105378.
- Hakimov, Rustamdjan, and Dorothea Kübler. 2021. "Experiments on Centralized School Choice and College Admissions: A Survey." *Experimental Econ.* 24:434488.
- Hakimov, Rustamdjan, Dorothea Kübler, and Siqi Pan. 2021. "Costly Information Acquisition in Centralized Matching Markets." Discussion Paper no. 280, Collaborative Res. Center Transregio 190, Berlin.
- Hassidim, Avinatan, Assaf Romm, and Ran I. Shorrer. 2021. "The Limits of Incentives in Economic Matching Procedures." *Management Sci.* 67 (2): 951–63.
- Hatfield, John William, Fuhito Kojima, and Scott Duke Kominers. 2014. "Investment Incentives in Labor Market Matching." *A.E.R.* 104 (5): 436–41.
- He, YingHua. 2017. "Gaming the Boston School Choice Mechanism in Beijing." Working paper, Toulouse School Econ.
- Immorlica, Nicole S., Jacob D. Leshno, Irene Y. Lo, and Brendan J. Lucier. 2020. "Information Acquisition in Matching Markets: The Role of Price Discovery." Working paper.
- Kahneman, Daniel. 2003. "Maps of Bounded Rationality: Psychology for Behavioral Economics." *A.E.R.* 93 (5): 1449–75.
- . 2011. *Thinking, Fast and Slow*. New York: Farrar, Straus & Giroux.
- Kahneman, Daniel, Jack L. Knetsch, and Richard H. Thaler. 1990. "Experimental Tests of the Endowment Effect and the Coase Theorem." *J.P.E.* 98 (6): 1325–48.

- Kapor, Adam J., Christopher A. Neilson, and Seth D. Zimmerman. 2020. "Heterogeneous Beliefs and School Choice Mechanisms." *A.E.R.* 110 (5): 1274–315.
- Ke, T. Tony, Zuo-Jun Max Shen, and J. Miguel Villas-Boas. 2016. "Search for Information on Multiple Products." *Management Sci.* 62 (12): 3576–603.
- Ke, T. Tony, and J. Miguel Villas-Boas. 2019. "Optimal Learning before Choice." *J. Econ. Theory* 180:383–437.
- Klijn, Flip, Joana Pais, and Marc Vorsatz. 2019. "Static versus Dynamic Deferred Acceptance in School Choice: Theory and Experiment." *Games and Econ. Behavior* 113:147–63.
- Loewenstein, George, and Ted O'Donoghue. 2004. "Animal Spirits: Affective and Deliberative Processes in Economic Behavior." Working Paper no. 04-14, Center Analytic Econ., Cornell Univ., Ithaca, NY.
- Manski, Charles F., and David A. Wise. 1983. *College Choice in America*. Cambridge, MA: Harvard Univ. Press.
- Narita, Yusuke. 2018. "Match or Mismatch? Learning and Inertia in School Choice." Working paper, Dept. Econ., Yale Univ., New Haven, CT.
- Niederle, Muriel, and Leeat Yariv. 2009. "Decentralized Matching with Aligned Preferences." Working Paper no. 14840, NBER, Cambridge, MA.
- Nöldeke, Georg, and Larry Samuelson. 2015. "Investment and Competitive Matching." *Econometrica* 83 (3): 835–96.
- Rees-Jones, Alex, and Samuel Skowronek. 2018. "An Experimental Investigation of Preference Misrepresentation in the Residency Match." *Proc. Nat. Acad. Sci. USA* 115 (45): 11471–76.
- Roth, Alvin E. 1990. "New Physicians: A Natural Experiment in Market Organization." *Science* 250:1524–28.
- Roth, Alvin E., and Elliott Peranson. 1999. "The Redesign of the Matching Market for American Physicians: Some Engineering Aspects of Economic Design." *A.E.R.* 89 (4): 748–80.
- Roth, Alvin E., and Uriel G. Rothblum. 1999. "Truncation Strategies in Matching Markets—In Search of Advice for Participants." *Econometrica* 67 (1): 21–43.
- Roth, Alvin E., and Marilda A. Oliveira Sotomayor. 1990. *Two-Sided Matching: A Study in Game-Theoretic Modeling and Analysis*. Econometric Soc. Monograph Series. Cambridge: Cambridge Univ. Press.
- Shorrer, Ran L., and Sándor Sóvágó. 2017. "Obvious Mistakes in a Strategically Simple College Admissions Environment." Discussion Paper no. 2017-107/V, Tinbergen Inst., Amsterdam.
- Weitzman, Martin L. 1979. "Optimal Search for the Best Alternative." *Econometrica* 47 (3): 641–54.
- Westkamp, Alexander. 2013. "An Analysis of the German University Admissions System." *Econ. Theory* 53 (3): 561–89.